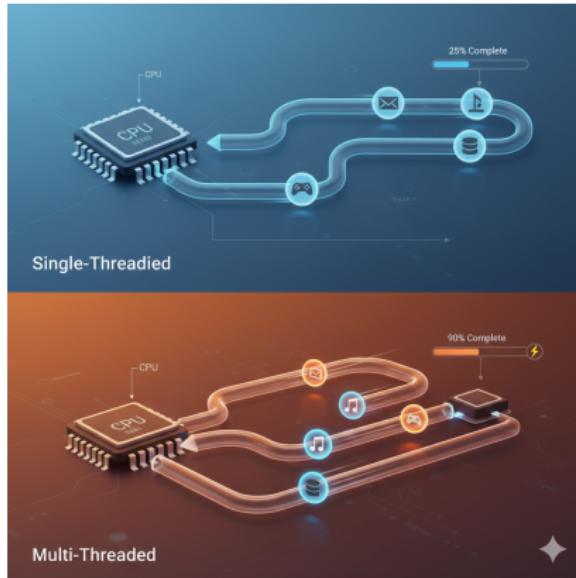


# CSL 301

## OPERATING SYSTEMS



### Lecture 16

#### Concurrency Intro Threads

Instructor  
Dr. Dhiman Saha

## Recap: The Process

- ▶ The OS gives us the abstraction of a **process**.

## Recap: The Process

- ▶ The OS gives us the abstraction of a **process**.
- ▶ A process has two key components:
  - ▶ **Virtual CPUs**: The illusion that the program is running on its own CPU.
  - ▶ **Virtual Memory**: The illusion that the program has its own private address space.

## Recap: The Process

- ▶ The OS gives us the abstraction of a **process**.
- ▶ A process has two key components:
  - ▶ **Virtual CPUs**: The illusion that the program is running on its own CPU.
  - ▶ **Virtual Memory**: The illusion that the program has its own private address space.
- ▶ This "classic" view of a process has a **single point of execution**.
  - ▶ One Program Counter (PC).
  - ▶ One set of instructions being executed.

## A New Abstraction: The Thread

- ▶ A **multi-threaded** program has *more than one* point of execution.
- ▶ Think of it as multiple PCs, each fetching and executing instructions.
- ▶ It's like having multiple processes, with one key difference...

# A New Abstraction: The Thread

- ▶ A **multi-threaded** program has *more than one* point of execution.
- ▶ Think of it as multiple PCs, each fetching and executing instructions.
- ▶ It's like having multiple processes, with one key difference...

## The Key Difference

All threads within a single process **share the same address space**. They can all access the same data.

# Thread vs. Process State

## What's shared?

- ▶ Address Space (Code, Heap)
- ▶ Page Table
- ▶ File Descriptors

## What's private?

- ▶ Program Counter (PC)
- ▶ Registers
- ▶ **Stack**

# Thread vs. Process State

## What's shared?

- ▶ Address Space (Code, Heap)
- ▶ Page Table
- ▶ File Descriptors

## What's private?

- ▶ Program Counter (PC)
- ▶ Registers
- ▶ Stack

- ▶ Switching between threads of the same process is much cheaper than switching between processes. Why?

# Thread vs. Process State

## What's shared?

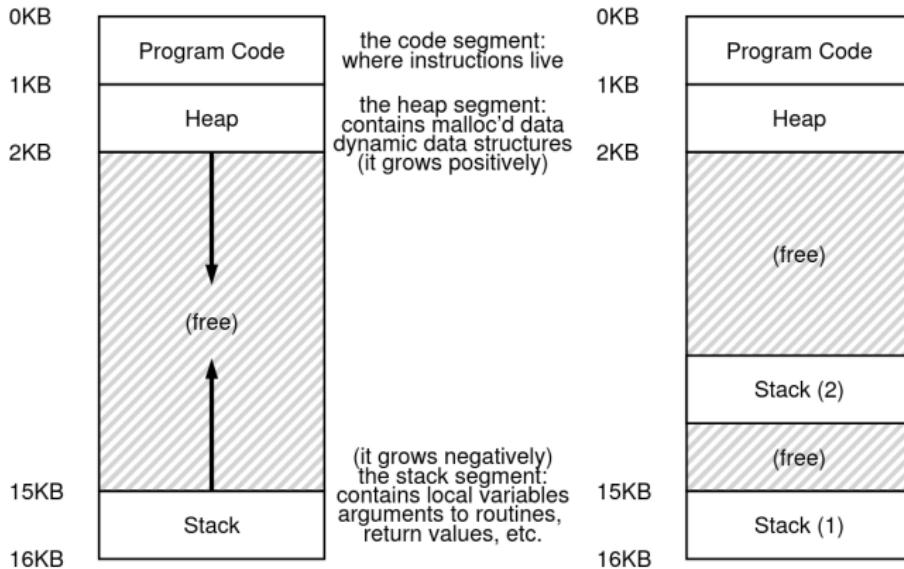
- ▶ Address Space (Code, Heap)
- ▶ Page Table
- ▶ File Descriptors

## What's private?

- ▶ Program Counter (PC)
- ▶ Registers
- ▶ **Stack**

- ▶ Switching between threads of the same process is much cheaper than switching between processes. Why?
- ▶ *No need to switch the address space!*

# Visualizing Address Spaces - Single Vs Multi-Threaded



- ▶ In a multi-threaded process, there is one stack **per thread**.
- ▶ This means local variables and function call arguments are private to each thread (this is called *thread-local storage*).

## Reason 1: Parallelism

- ▶ Imagine a program performing an operation on a very large array.
- ▶ On a system with multiple CPUs, you can speed this up by dividing the work.

## Reason 1: Parallelism

- ▶ Imagine a program performing an operation on a very large array.
- ▶ On a system with multiple CPUs, you can speed this up by dividing the work.
- ▶ **Parallelization:** The task of transforming a single-threaded program into one that does work on multiple CPUs.
- ▶ Using one thread per CPU is a natural way to make programs run faster on modern hardware.

## Reason 2: Overlapping I/O

- ▶ Many programs perform slow I/O operations (e.g., reading from a disk, waiting for a network request).
- ▶ In a single-threaded program, the entire process **blocks** and can't do any other work.

## Reason 2: Overlapping I/O

- ▶ Many programs perform slow I/O operations (e.g., reading from a disk, waiting for a network request).
- ▶ In a single-threaded program, the entire process **blocks** and can't do any other work.
- ▶ With threads, you can avoid getting stuck!
  - ▶ While one thread is blocked waiting for I/O...
  - ▶ ...the OS scheduler can switch to another thread, which is ready to run and do useful work.

## Reason 2: Overlapping I/O

- ▶ Many programs perform slow I/O operations (e.g., reading from a disk, waiting for a network request).
- ▶ In a single-threaded program, the entire process **blocks** and can't do any other work.
- ▶ With threads, you can avoid getting stuck!
  - ▶ While one thread is blocked waiting for I/O...
  - ▶ ...the OS scheduler can switch to another thread, which is ready to run and do useful work.
- ▶ This is essential for modern servers (web servers, databases) that handle many concurrent requests.

## Example: Creating Threads (t0.c)

```
1 #include <stdio.h>
2 #include <pthread.h>
3
4 void *mythread(void *arg) {
5     printf("%s\n", (char *) arg);
6     return NULL;
7 }
8
9 int main(int argc, char *argv[]) {
10    pthread_t p1, p2;
11    printf("main: begin\n");
12    pthread_create(&p1, NULL, mythread, "A");
13    pthread_create(&p2, NULL, mythread, "B");
14    // join waits for the threads to finish
15    pthread_join(p1, NULL);
16    pthread_join(p2, NULL);
17    printf("main: end\n");
18    return 0;
19 }
```

## Understanding Thread Execution

- ▶ The main thread creates two new threads, T1 and T2.
- ▶ It then calls `pthread_join()` twice, waiting for each thread to complete.
- ▶ What will the output be?

# Understanding Thread Execution

- ▶ The main thread creates two new threads, T1 and T2.
- ▶ It then calls `pthread_join()` twice, waiting for each thread to complete.
- ▶ What will the output be?

**Trace 1**

```
1 main: begin
2 A
3 B
4 main: end
```

**Trace 2**

```
1 main: begin
2 B
3 A
4 main: end
```

**Trace 3**

```
1 main: begin
2 main: end
3 A
4 B
```

# Understanding Thread Execution

- ▶ The main thread creates two new threads, T1 and T2.
- ▶ It then calls `pthread_join()` twice, waiting for each thread to complete.
- ▶ What will the output be?

Trace 1

```
1 main: begin  
2 A  
3 B  
4 main: end
```

Trace 2

```
1 main: begin  
2 B  
3 A  
4 main: end
```

Trace 3

```
1 main: begin  
2 main: end  
3 A  
4 B
```

## The Point

The order of execution is non-deterministic! It depends on the OS scheduler. Any of these outputs (and more) are possible.

## What if Threads Interact?

- ▶ The previous example was simple: the threads didn't interact.
- ▶ Things get much more complicated when threads access **shared data**.
- ▶ Let's look at an example where two threads try to update a shared counter.

## Example: A Shared Counter (t1.c) |

```
1 #include <stdio.h>
2 #include <pthread.h>
3
4 static volatile int counter = 0;
5
6 void *mythread(void *arg) {
7     printf("%s: begin\n", (char *) arg);
8     int i;
9     for (i = 0; i < 1e7; i++) {
10         counter = counter + 1;
11     }
12     printf("%s: done\n", (char *) arg);
13     return NULL;
14 }
```

## Example: A Shared Counter (t1.c) ||

```
19 | int main(int argc, char *argv[]) {
20 |     pthread_t p1, p2;
21 |     pthread_create(&p1, NULL, mythread, "A");
22 |     pthread_create(&p2, NULL, mythread, "B");
23 |     pthread_join(p1, NULL);
24 |     pthread_join(p2, NULL);
25 |     printf("main: done with both (counter = %d)\\n", counter);
26 |     return 0;
27 | }
```

## What Should The Result Be?

- ▶ Two threads, each incrementing the counter 10,000,000 times.
- ▶ Expected final value: **20,000,000**.

# What Should The Result Be?

- ▶ Two threads, each incrementing the counter 10,000,000 times.
- ▶ Expected final value: **20,000,000**.
- ▶ But when we run it...

```
1 prompt> ./main
2 main: begin (counter = 0)
3 A: begin
4 B: begin
5 A: done
6 B: done
7 main: done with both (counter = 19345221)
8
9 prompt> ./main
10 main: done with both (counter = 19221041)
```

## Problem!

Not only is the result wrong, it's **different** every time! Why?

## The "Atomic" Illusion

- ▶ The C statement `counter = counter + 1;` seems simple.
- ▶ But to the CPU, it's not a single, **atomic** operation.
- ▶ It's actually a sequence of instructions.

# The "Atomic" Illusion

- ▶ The C statement `counter = counter + 1;` seems simple.
- ▶ But to the CPU, it's not a single, **atomic** operation.
- ▶ It's actually a sequence of instructions.

## x86 Assembly for 'counter++'

```
1 mov 0x8049a1c, %eax      ; Load counter's value
2 add $0x1, %eax          ; Add 1 to register
3 mov %eax, 0x8049a1c     ; Store new value
```

# The "Atomic" Illusion

- ▶ The C statement `counter = counter + 1;` seems simple.
- ▶ But to the CPU, it's not a single, **atomic** operation.
- ▶ It's actually a sequence of instructions.

## x86 Assembly for 'counter++'

```
1 mov 0x8049a1c, %eax      ; Load counter's value
2 add $0x1, %eax          ; Add 1 to register
3 mov %eax, 0x8049a1c     ; Store new value
```

## The Core Problem

A context switch can happen **between any** of these instructions!

# Anatomy of a Race Condition

Let's say counter is 50.

- ▶ **Thread 1** executes `mov`, loading 50 into its private register (`eax=50`).

# Anatomy of a Race Condition

Let's say counter is 50.

- ▶ **Thread 1** executes `mov`, loading 50 into its private register (`eax=50`).
- ▶ **Thread 1** executes `add`, incrementing its register (`eax=51`).

## The Result

An increment has been lost! We performed two increments, but the value only increased by one.

# Anatomy of a Race Condition

Let's say counter is 50.

- ▶ **Thread 1** executes `mov`, loading 50 into its private register (`eax=50`).
- ▶ **Thread 1** executes `add`, incrementing its register (`eax=51`).
- ▶ **TIMER INTERRUPT!** The OS context switches to Thread 2.

## The Result

An increment has been lost! We performed two increments, but the value only increased by one.

# Anatomy of a Race Condition

Let's say counter is 50.

- ▶ **Thread 1** executes `mov`, loading 50 into its private register (`eax=50`).
- ▶ **Thread 1** executes `add`, incrementing its register (`eax=51`).
- ▶ **TIMER INTERRUPT!** The OS context switches to Thread 2.
- ▶ **Thread 2** executes `mov`, loading the *original* value 50 from memory into its private register (`eax=50`).

## The Result

An increment has been lost! We performed two increments, but the value only increased by one.

# Anatomy of a Race Condition

Let's say counter is 50.

- ▶ **Thread 1** executes `mov`, loading 50 into its private register (`eax=50`).
- ▶ **Thread 1** executes `add`, incrementing its register (`eax=51`).
- ▶ **TIMER INTERRUPT!** The OS context switches to Thread 2.
- ▶ **Thread 2** executes `mov`, loading the *original* value 50 from memory into its private register (`eax=50`).
- ▶ **Thread 2** executes `add` (`eax=51`) and then `mov`, storing 51 back to memory. counter is now 51.

## The Result

An increment has been lost! We performed two increments, but the value only increased by one.

# Anatomy of a Race Condition

Let's say counter is 50.

- ▶ **Thread 1** executes `mov`, loading 50 into its private register (`eax=50`).
- ▶ **Thread 1** executes `add`, incrementing its register (`eax=51`).
- ▶ **TIMER INTERRUPT!** The OS context switches to Thread 2.
- ▶ **Thread 2** executes `mov`, loading the *original* value 50 from memory into its private register (`eax=50`).
- ▶ **Thread 2** executes `add` (`eax=51`) and then `mov`, storing 51 back to memory. counter is now 51.
- ▶ **CONTEXT SWITCH!** The OS switches back to Thread 1.

## The Result

An increment has been lost! We performed two increments, but the value only increased by one.

# Anatomy of a Race Condition

Let's say counter is 50.

- ▶ **Thread 1** executes `mov`, loading 50 into its private register (`eax=50`).
- ▶ **Thread 1** executes `add`, incrementing its register (`eax=51`).
- ▶ **TIMER INTERRUPT!** The OS context switches to Thread 2.
- ▶ **Thread 2** executes `mov`, loading the *original* value 50 from memory into its private register (`eax=50`).
- ▶ **Thread 2** executes `add` (`eax=51`) and then `mov`, storing 51 back to memory. counter is now 51.
- ▶ **CONTEXT SWITCH!** The OS switches back to Thread 1.
- ▶ **Thread 1** resumes. It executes its final instruction, `mov`, storing its register value (51) back to memory. counter is still 51.

## The Result

An increment has been lost! We performed two increments, but the value only increased by one.

# Key Concurrency Terms

**Race Condition** The result of a program depends on the timing or interleaving of operations. The situation we just saw.

# Key Concurrency Terms

**Race Condition** The result of a program depends on the timing or interleaving of operations. The situation we just saw.

**Critical Section** A piece of code that accesses a shared resource (like our counter) and must not be concurrently executed by more than one thread.

# Key Concurrency Terms

**Race Condition** The result of a program depends on the timing or interleaving of operations. The situation we just saw.

**Critical Section** A piece of code that accesses a shared resource (like our counter) and must not be concurrently executed by more than one thread.

**Mutual Exclusion** The property we want to enforce. It guarantees that if one thread is executing in a critical section, other threads will be prevented from entering it.

# The Wish for Atomicity

- ▶ If we had a single hardware instruction that did the load, add, and store all at once, there would be no problem.

```
1 memory-add 0x8049a1c , $0x1
```

- ▶ This is an **atomic** operation: it runs "as a unit", cannot be interrupted, and appears to happen instantaneously.

# The Wish for Atomicity

- ▶ If we had a single hardware instruction that did the load, add, and store all at once, there would be no problem.

```
1 memory-add 0x8049a1c , $0x1
```

- ▶ This is an **atomic** operation: it runs "as a unit", cannot be interrupted, and appears to happen instantaneously.
- ▶ But hardware can't provide atomic instructions for every complex operation we might want (e.g., "atomically update B-Tree").

## The Real Solution: Synchronization

- ▶ Instead of asking for complex atomic instructions...
- ▶ We ask the hardware for a few simple, useful atomic instructions.
- ▶ We then use these, with help from the OS, to build higher-level **synchronization primitives**.
- ▶ Examples: Locks (Mutexes), Semaphores, Condition Variables.

# The Real Solution: Synchronization

- ▶ Instead of asking for complex atomic instructions...
- ▶ We ask the hardware for a few simple, useful atomic instructions.
- ▶ We then use these, with help from the OS, to build higher-level **synchronization primitives**.
- ▶ Examples: Locks (Mutexes), Semaphores, Condition Variables.
- ▶ By using these primitives, we can protect our critical sections and ensure mutual exclusion.

## The Goal

To build multi-threaded code that is correct, reliable, and produces deterministic results despite the challenging nature of concurrency.