

# CSL 301

## OPERATING SYSTEMS



## Lecture 28

### Authentication in Operating Systems

Instructor  
Dr. Dhiman Saha

Image Source: Gemini

# The Need for Authentication

- ▶ Operating systems provide services, and many have security implications.
- ▶ The OS must decide whether to perform an action or deny it based on security policy.
- ▶ **Context is everything!** The most important context is *who* is asking.

# The Need for Authentication

- ▶ Operating systems provide services, and many have security implications.
- ▶ The OS must decide whether to perform an action or deny it based on security policy.
- ▶ **Context is everything!** The most important context is *who* is asking.

## Key Terminology

- ▶ **Principal:** The entity (e.g., human user, group) making a request.
- ▶ **Agent:** A process executing on behalf of a principal.
- ▶ **Object:** The resource being accessed (e.g., file, network port).

# The Need for Authentication

- ▶ Operating systems provide services, and many have security implications.
- ▶ The OS must decide whether to perform an action or deny it based on security policy.
- ▶ **Context is everything!** The most important context is *who* is asking.

## Key Terminology

- ▶ **Principal:** The entity (e.g., human user, group) making a request.
- ▶ **Agent:** A process executing on behalf of a principal.
- ▶ **Object:** The resource being accessed (e.g., file, network port).

## The Fundamental Problem

How does the operating system reliably determine the identity of the principal behind a request to enforce security goals?

# Attaching Identities to Processes

- ▶ Processes are the actors in an OS. Everything is done by a process.
- ▶ If we can securely identify processes, we can enforce security policies.
- ▶ **Identity Inheritance:** Typically, a new process (child) inherits the identity of its creator (parent).

# Attaching Identities to Processes

- ▶ Processes are the actors in an OS. Everything is done by a process.
- ▶ If we can securely identify processes, we can enforce security policies.
- ▶ **Identity Inheritance:** Typically, a new process (child) inherits the identity of its creator (parent).

## The Starting Point: User Login

How does the very first process for a user (e.g., the command shell) get its identity?

- ▶ The user must first **log in**.
- ▶ The OS must be able to verify the identity of the human user.
- ▶ This verified identity is then attached to the user's initial process.

## Crux

The entire security of the system relies on the OS making the right decision at the authentication stage.

# The Three Factors of Authentication

Classically, a user's identity can be verified based on:

## 1. What You Know

A secret, like a password or a PIN.

## 2. What You Have

A physical object, like a key card, smartphone, or USB token.

## 3. What You Are

A biological characteristic (biometrics), like a fingerprint or retina scan.

## Factor 1: What You Know (Passwords)

- ▶ The most common form of authentication.
- ▶ A secret shared between the user and the system.

## Factor 1: What You Know (Passwords)

- ▶ The most common form of authentication.
- ▶ A secret shared between the user and the system.

### Major Security Flaw: How to Store the Secret?

If the system stores the password in plaintext, anyone who accesses the password file learns all user passwords. This is a huge vulnerability.

## Factor 1: What You Know (Passwords)

- ▶ The most common form of authentication.
- ▶ A secret shared between the user and the system.

### Major Security Flaw: How to Store the Secret?

If the system stores the password in plaintext, anyone who accesses the password file learns all user passwords. This is a huge vulnerability.

### Solution: Store Hashes, Not Passwords

- ▶ A **hash function** is a one-way function: easy to compute, but infeasible to reverse.
- ▶ **Storage:** The system stores ‘hash (password)’.
- ▶ **Verification:** When a user provides a password, the system computes its hash and compares it with the stored hash.
- ▶ If an attacker steals the password file, they only get the hashes, not the passwords themselves.

# Attacking Hashed Passwords

Even with hashing, passwords can be weak.

## Problem 1: Guessing Attacks

- ▶ **Brute-force:** Trying all possible combinations. Infeasible for long, complex passwords.
- ▶ **Dictionary Attack:** Trying common words, names, and simple patterns (e.g., "password123"). Surprisingly effective as humans choose memorable (and thus predictable) passwords.

# Attacking Hashed Passwords

Even with hashing, passwords can be weak.

## Problem 1: Guessing Attacks

- ▶ **Brute-force:** Trying all possible combinations. Infeasible for long, complex passwords.
- ▶ **Dictionary Attack:** Trying common words, names, and simple patterns (e.g., "password123"). Surprisingly effective as humans choose memorable (and thus predictable) passwords.

## Problem 2: Pre-computed Hashes (Rainbow Tables)

- ▶ If an attacker steals the password file, they can guess passwords offline.
- ▶ An attacker can pre-compute the hashes of a large dictionary of words.
- ▶ Then simply compare the stolen hashes against their pre-computed list to find matches instantly.

# Defense: Password Salting

## The Problem

With a standard hash, the same password always produces the same hash, making pre-computation attacks possible.

# Defense: Password Salting

## The Problem

With a standard hash, the same password always produces the same hash, making pre-computation attacks possible.

## Solution: Use a "Salt"

- ▶ A **salt** is a unique, random value generated for each user.
- ▶ **Storage:** The system stores the user's salt (in plaintext) and ' $\text{hash}(\text{password} + \text{salt})$ '.
- ▶ **Verification:** The system retrieves the user's salt, concatenates it with the provided password, and hashes the result for comparison.

# Defense: Password Salting

## The Problem

With a standard hash, the same password always produces the same hash, making pre-computation attacks possible.

## Solution: Use a "Salt"

- ▶ A **salt** is a unique, random value generated for each user.
- ▶ **Storage:** The system stores the user's salt (in plaintext) and ' $\text{hash}(\text{password} + \text{salt})$ '.
- ▶ **Verification:** The system retrieves the user's salt, concatenates it with the provided password, and hashes the result for comparison.

**How does this help?** An attacker must now compute hashes for every possible salt. If the salt is large enough, pre-computing a rainbow table becomes infeasible. The dictionary attack must be run separately for each stolen password, using its unique salt.

## Factor 2: What You Have (Tokens)

Authentication based on possessing a physical object.

- ▶ **Examples:** ATM cards, smart cards, USB security tokens (e.g., YubiKey), smartphones.
- ▶ Often use a changing piece of information (like a one-time password or a cryptographic challenge-response) to prevent an attacker from simply copying the secret.
- ▶ This converts a static "what you know" (the secret on the device) into a dynamic "what you have".

## Factor 2: What You Have (Tokens)

Authentication based on possessing a physical object.

- ▶ **Examples:** ATM cards, smart cards, USB security tokens (e.g., YubiKey), smartphones.
- ▶ Often use a changing piece of information (like a one-time password or a cryptographic challenge-response) to prevent an attacker from simply copying the secret.
- ▶ This converts a static "what you know" (the secret on the device) into a dynamic "what you have".

### Primary Weakness

What if you lose the device?

- ▶ You are locked out.
- ▶ An attacker who finds it might be able to impersonate you.

## Factor 3: What You Are (Biometrics)

Authentication based on a unique physical or behavioral trait.

- ▶ **Examples:** Fingerprint, face recognition, iris/retina scan, voice print, typing pattern.

## Factor 3: What You Are (Biometrics)

Authentication based on a unique physical or behavioral trait.

- ▶ **Examples:** Fingerprint, face recognition, iris/retina scan, voice print, typing pattern.

### The Challenge: Imprecision

Biometric measurements are never exactly the same twice.

- ▶ A simple bit-for-bit comparison (like with a password hash) will fail.
- ▶ The system must allow for some "tolerance" or sloppiness.

## Factor 3: What You Are (Biometrics)

This leads to two types of errors:

- ▶ **False Negative:** The system fails to recognize a legitimate user. (Annoying for the user).
- ▶ **False Positive:** The system incorrectly authenticates an imposter. (Disastrous for security).

### Crossover Error Rate (CER)

There is an inherent trade-off. Tuning the system to be more secure (lower false positives) often makes it less convenient (higher false negatives). The CER is a metric for the accuracy of a biometric system.

# Multi-Factor Authentication (MFA)

Security is significantly increased by requiring evidence from **more than one** authentication factor.

## Example: ATM Transaction

1. **What you have:** The physical ATM card.
2. **What you know:** The Personal Identification Number (PIN).

# Multi-Factor Authentication (MFA)

Security is significantly increased by requiring evidence from **more than one** authentication factor.

## Example: ATM Transaction

1. **What you have:** The physical ATM card.
2. **What you know:** The Personal Identification Number (PIN).

## Example: Modern Web Login

1. **What you know:** Your password.
2. **What you have:** A one-time code sent to your smartphone.

# Multi-Factor Authentication (MFA)

Security is significantly increased by requiring evidence from **more than one** authentication factor.

## Example: ATM Transaction

1. **What you have:** The physical ATM card.
2. **What you know:** The Personal Identification Number (PIN).

## Example: Modern Web Login

1. **What you know:** Your password.
2. **What you have:** A one-time code sent to your smartphone.

## Benefit of MFA

An attacker must compromise multiple, independent factors.

Stealing your password is not enough if they don't also have your phone.

# Authenticating Non-Humans

- ▶ Not all processes are run by a human user (e.g., a web server, a database service).
- ▶ These services still need an identity to manage their privileges.

# Authenticating Non-Humans

- ▶ Not all processes are run by a human user (e.g., a web server, a database service).
- ▶ These services still need an identity to manage their privileges.

## Solution

- ▶ Create a dedicated "user account" for the service (e.g., a user named 'webserver').
- ▶ This account has only the privileges necessary for its task (Principle of Least Privilege).

# Authenticating Non-Humans

## How does the service get this identity?

- ▶ A system administrator logs in and starts the service.
- ▶ Privileged commands like `sudo` can be used to launch a process under a different user identity.

```
sudo -u webserver apache2
```

- ▶ The ‘apache2’ process (and any children it creates) will run with the ‘webserver’ user’s identity and privileges.

## Example: Linux Login Procedure

1. A privileged ‘login’ process prompts for a username.
2. The process then prompts for the password (which is not echoed to the screen).
3. It looks up the username in the system’s password database (e.g., /etc/shadow).
  - ▶ If the user is not found, the login fails.
  - ▶ If found, it retrieves the user’s **salt** and stored **password hash**.
4. The ‘login’ process hashes the entered password with the retrieved salt and compares it to the stored hash.
  - ▶ If they do not match, the login fails.
5. **If they match:**
  - ▶ It calls `fork()` to create a new process.
  - ▶ This new process’s user and group ID are set to that of the authenticated user.
  - ▶ It then calls `exec()` to replace itself with the user’s default shell (e.g., /bin/bash), which now runs with the user’s identity.

## Summary

- ▶ Authentication is the bedrock of OS security, used to verify the identity of a principal.
- ▶ The process begins with a user login, which attaches a verified identity to an initial process. This identity is then inherited by child processes.
- ▶ Authentication relies on three factors: something you **know**, **have**, or **are**.
- ▶ Passwords (*what you know*) must be stored securely using **hashes** and **salts** to protect against theft and guessing attacks.
- ▶ Biometrics (*what you are*) are powerful but must manage the trade-off between security (false positives) and convenience (false negatives).
- ▶ **Multi-Factor Authentication (MFA)** provides the strongest security by combining two or more factors.