

CSL 301

OPERATING SYSTEMS

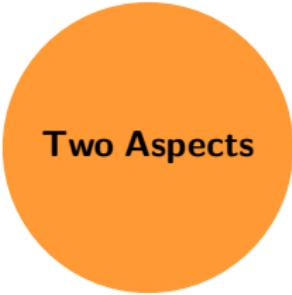
RESTRICTED

Lecture 3

Limited Direct Execution

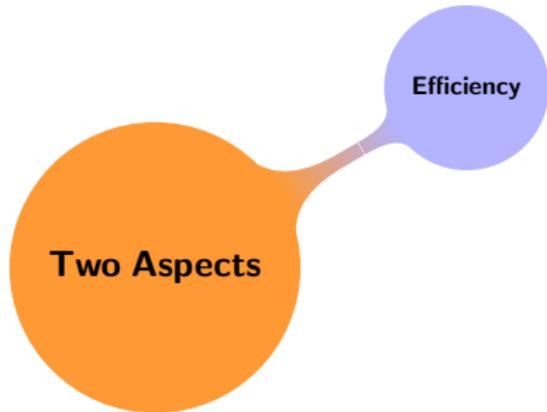
Instructor
Dr. Dhiman Saha

How to efficiently virtualize the CPU with control?

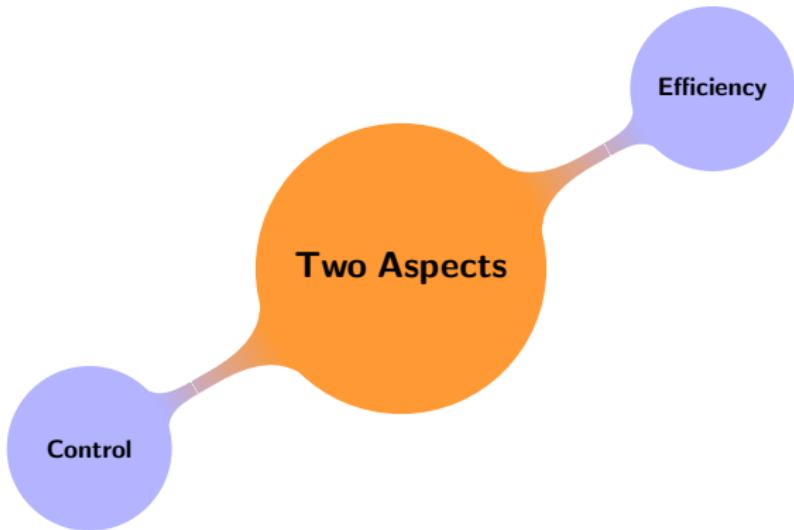


Two Aspects

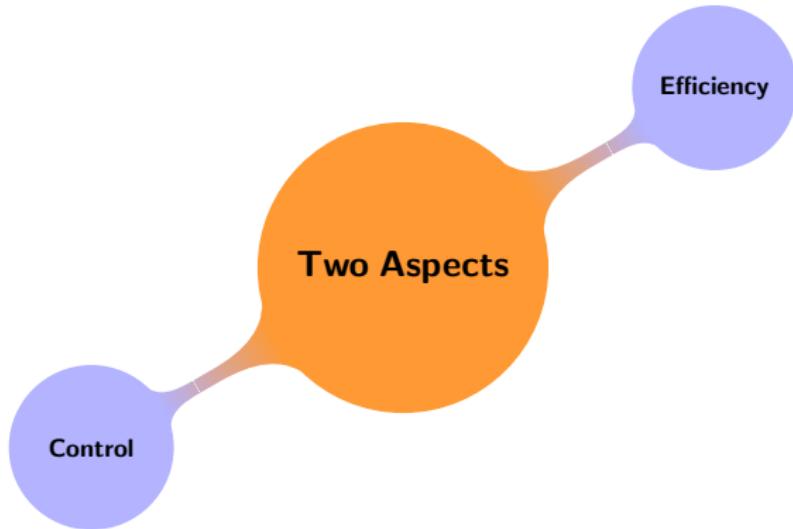
How to efficiently virtualize the CPU with control?



How to efficiently virtualize the CPU with control?



How to efficiently virtualize the CPU with control?



Basic Idea

Direct Execution

Direct Execution Protocol (Without Limits)

OS	Program
Create entry for process list	
Allocate memory for program	
Load program into memory	
Set up stack with argc/argv	
Clear registers	
Execute call main()	Run main() Execute return from main
Free memory of process	
Remove from process list	

- ▶ If we just run a program, how can the OS make sure the program doesn't do anything that we don't want it to do, while still running it efficiently?
- ▶ When we are running a process, how does the operating system stop it from running and switch to another process?

“Limited” Direct Execution

- ▶ A process must be able to perform I/O and some other restricted operations, but without giving the process complete control over the system.
- ▶ How can the OS and hardware work together to do so?

What if?

We let any process do whatever it wants in terms of I/O and other related operations.

- ▶ A new processor mode

user mode

Code that runs in user mode is **restricted** in what it can do.

Example

- ▶ In user mode, a process can't issue I/O requests
- ▶ Doing so would result in the processor raising an **exception**
- ▶ The OS would then likely kill the process.

- ▶ The OS (or kernel) run in this mode

In this mode, code that runs can do what it likes,

- ▶ including privileged operations
- ▶ issuing I/O requests
- ▶ executing all types of restricted instructions.

user mode \leftrightarrow kernel mode

system call

What happens in a simple function call?

- ▶ function call \implies a jump instruction
- ▶ stack \leftarrow new stack frame pushed
- ▶ Stack Pointer (SP) \leftarrow updated
- ▶ Old value of PC (return value) pushed to stack and PC updated
- ▶ Once function finishes, stack is popped and old PC value is retrieved effecting a return to the calling function

Stack frame

Contains return value, function arguments etc.

What changes in a system call?

- ▶ Switching from user to kernel mode
- ▶ The **trap** instruction (hidden from user)
- ▶ A separate **kernel stack** is used when in kernel mode. Why?
- ▶ Servicing the call. How?
- ▶ Switching from kernel to user mode
- ▶ The **return-from-trap** instruction

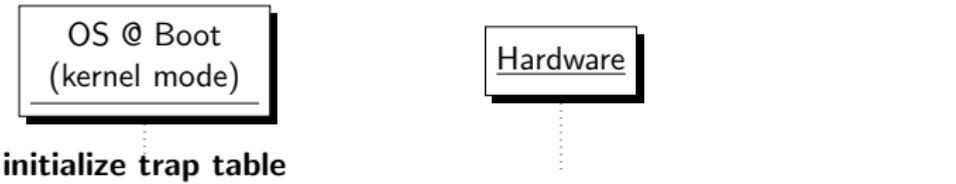
How does the **trap** know which code to run inside the OS?

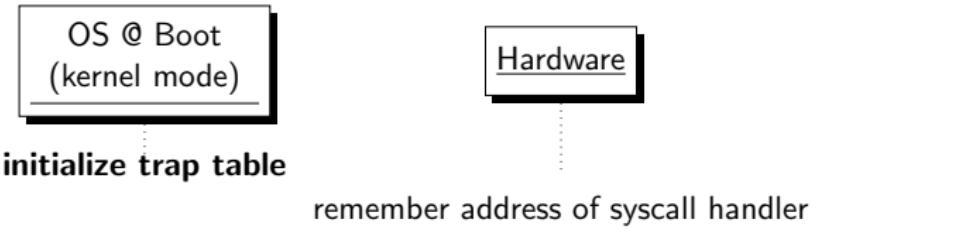
- ▶ Can/Should the calling process do this?
- ▶ The **trap table**
- ▶ Setup at **boot time** by the kernel

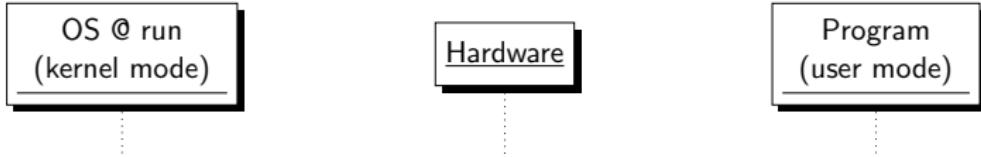
Limited Direct Execution Protocol

OS @ boot (kernel mode)	Hardware	
initialize trap table	remember address of... syscall handler	
OS @ run (kernel mode)	Hardware	Program (user mode)
Create entry for process list Allocate memory for program Load program into memory Setup user stack with argv Fill kernel stack with reg/PC return-from-trap	restore regs from kernel stack move to user mode jump to main	Run main() ... Call system call trap into OS
Handle trap Do work of syscall return-from-trap	save regs to kernel stack move to kernel mode jump to trap handler	
Free memory of process Remove from process list	restore regs from kernel stack move to user mode jump to PC after trap	... return from main trap (via <code>exit()</code>)











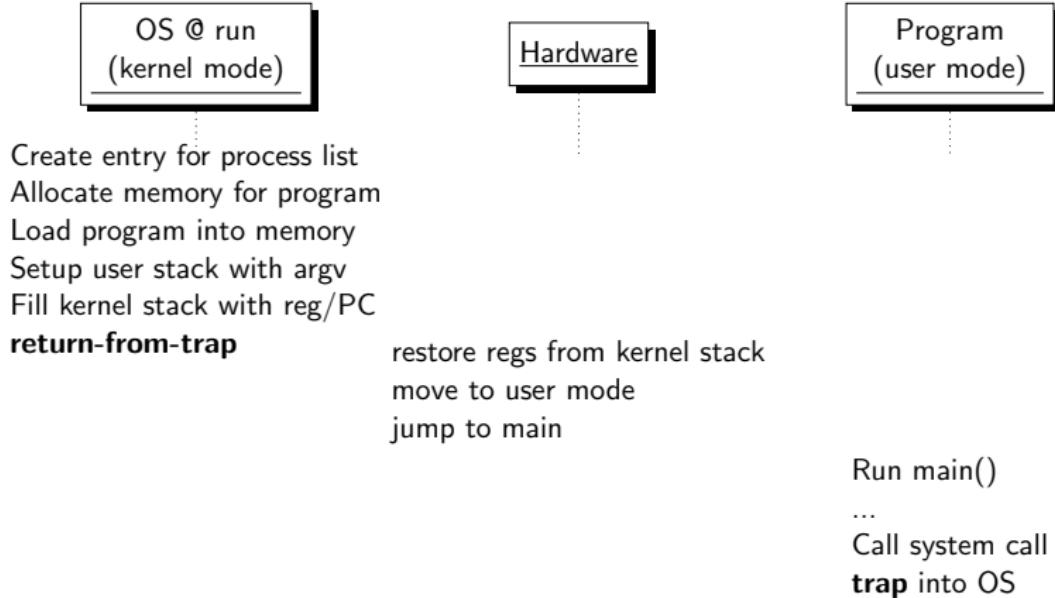
Create entry for process list
Allocate memory for program
Load program into memory
Setup user stack with argv
Fill kernel stack with reg/PC
return-from-trap

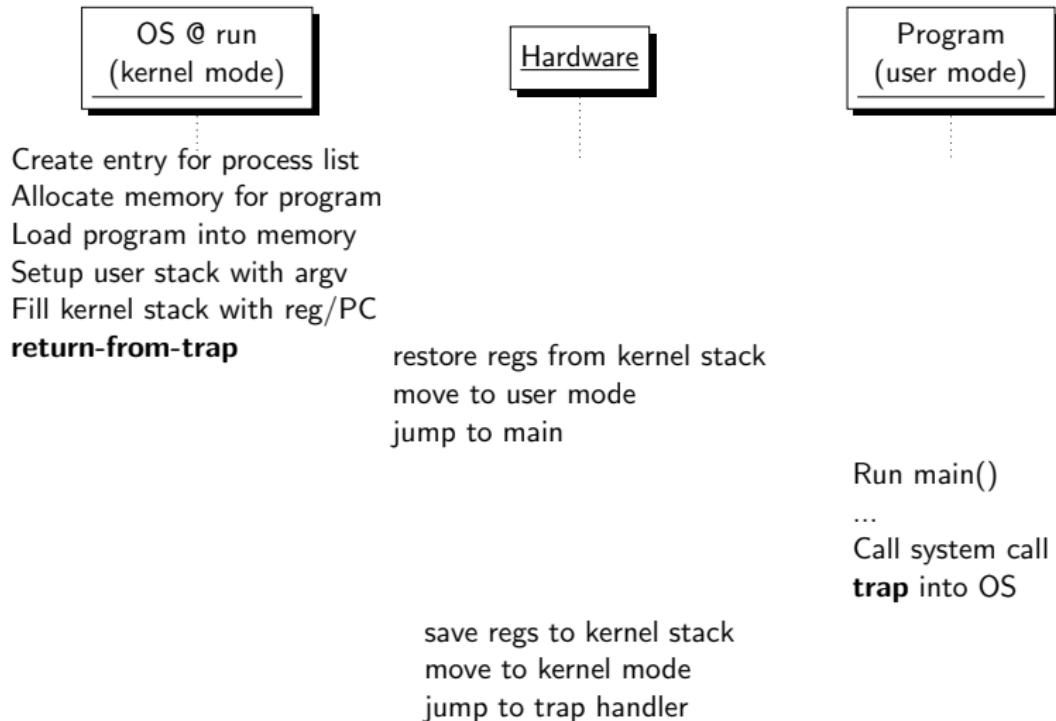


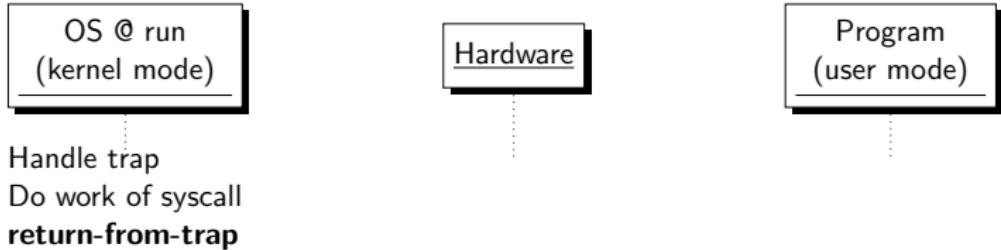
Create entry for process list
Allocate memory for program
Load program into memory
Setup user stack with argv
Fill kernel stack with reg/PC

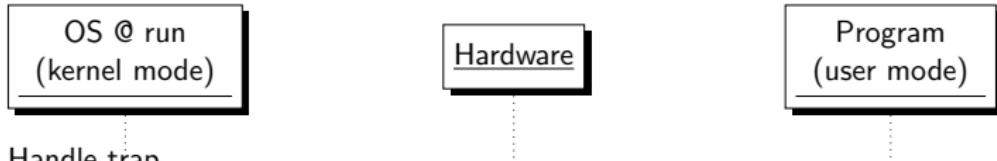
return-from-trap

restore regs from kernel stack
move to user mode
jump to main



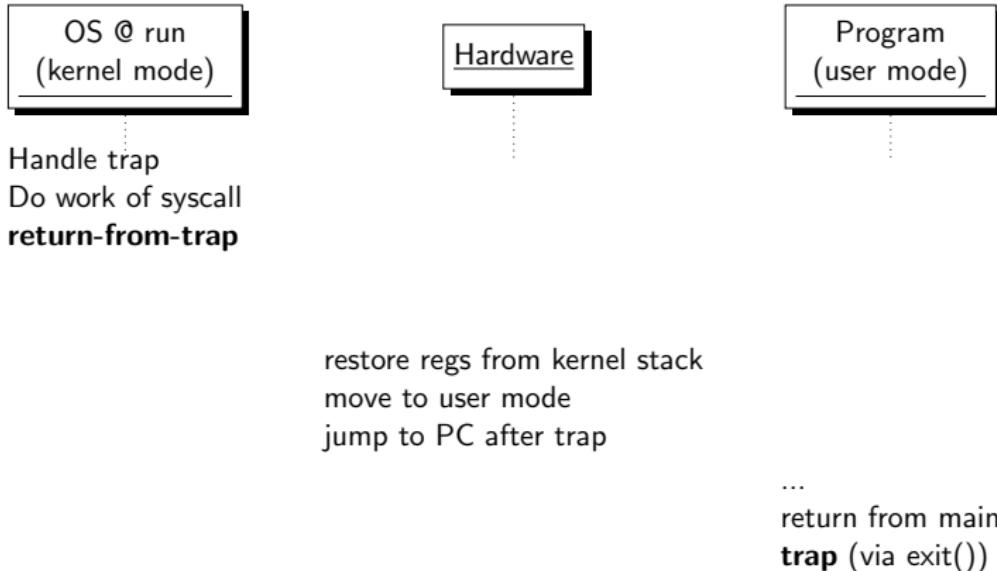


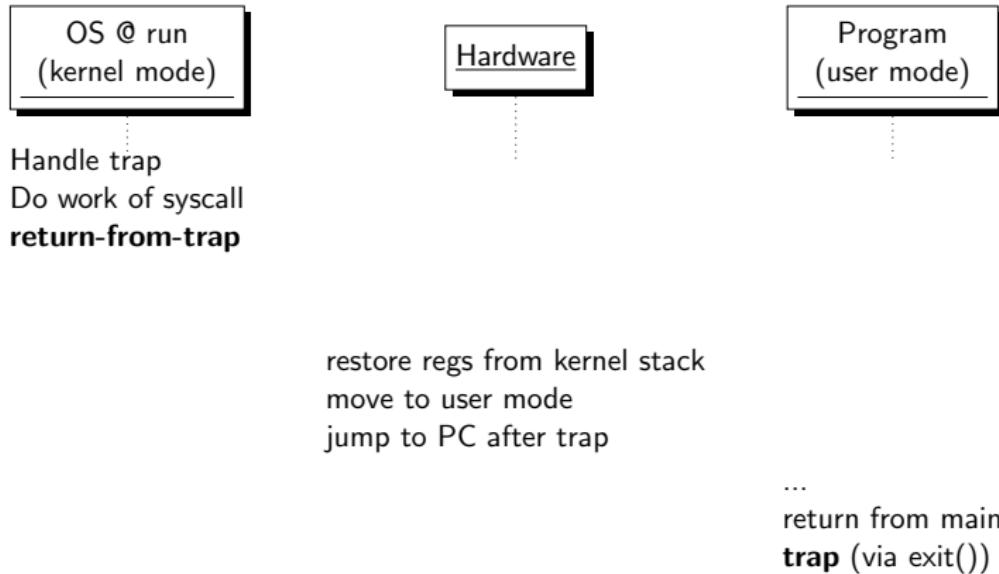




Handle trap
Do work of syscall
return-from-trap

restore regs from kernel stack
move to user mode
jump to PC after trap





Free memory of process
Remove from process list

- ▶ To specify the exact system call, a **system-call number** is usually assigned to each system call.
- ▶ The user code is thus responsible for placing the desired system-call number in a register or at a specified location on the stack
- ▶ The OS, when handling the system call inside the trap handler, examines this number, ensures it is valid, and, if it is, executes the corresponding code.

This gives a level of protection. How?

How can the operating system regain control of the CPU so that it can switch between processes?

Wait For System Calls

- ▶ Switch contexts for `syscall` interrupt.
- ▶ Provide special `yield()` system call.
- ▶ Applications also transfer control to the OS when they do something illegal.

In a cooperative scheduling system, the OS regains control of the CPU by waiting for a system call or an illegal operation of some kind to take place.

The OS Takes Control

- ▶ How can the OS gain control of the CPU even if processes are not being cooperative?
- ▶ What can the OS do to ensure a rogue process does not take over the machine?

A timer interrupt

When the interrupt is raised, the currently running process is halted, and a pre-configured interrupt handler in the OS runs.

Saving and Restoring Context

- ▶ All the OS has to do is save a few register values for the currently-executing process (onto its kernel stack)
- ▶ And restore a few for the soon-to-be-executing process (from its kernel stack).

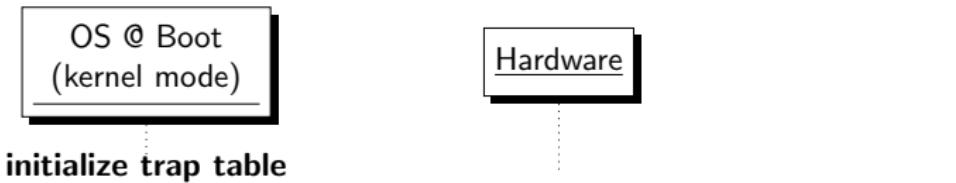
The xv6 Context Switch Code

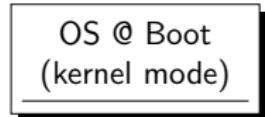
```
1 # void swtch(struct context **old, struct context *new);  
2 #  
3 # Save current register context in old  
4 # and then load register context from new.  
5 .globl swtch  
6 swtch:  
7     # Save old registers  
8     movl 4(%esp), %eax    # put old ptr into eax  
9     popl 0(%eax)          # save the old IP  
10    movl %esp, 4(%eax)    # and stack  
11    movl %ebx, 8(%eax)    # and other registers  
12    movl %ecx, 12(%eax)  
13    movl %edx, 16(%eax)  
14    movl %esi, 20(%eax)  
15    movl %edi, 24(%eax)  
16    movl %ebp, 28(%eax)  
17  
18     # Load new registers  
19     movl 4(%esp), %eax    # put new ptr into eax  
20     movl 28(%eax), %ebp    # restore other registers  
21     movl 24(%eax), %edi  
22     movl 20(%eax), %esi  
23     movl 16(%eax), %edx  
24     movl 12(%eax), %ecx  
25     movl 8(%eax), %ebx  
26     movl 4(%eax), %esp    # stack is switched here  
27     pushl 0(%eax)          # return addr put in place  
28     ret                  # finally return into new ctxt
```

LDE Protocol (Timer Interrupt)

OS @ boot (kernel mode)	Hardware	
initialize trap table	remember addresses of... syscall handler timer handler	
start interrupt timer	start timer interrupt CPU in X ms	
OS @ run (kernel mode)	Hardware	Program (user mode)
		Process A
		...
	timer interrupt save regs(A) to k-stack(A) move to kernel mode jump to trap handler	
Handle the trap Call switch() routine save regs(A) to proc-struct(A) restore regs(B) from proc-struct(B) switch to k-stack(B) return-from-trap (into B)		
	restore regs(B) from k-stack(B) move to user mode jump to B's PC	Process B
		...



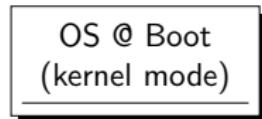




initialize trap table



remember addresses of...
syscall handler
timer handler

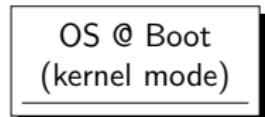


initialize trap table



remember addresses of...
syscall handler
timer handler

start interrupt timer



initialize trap table



remember addresses of...

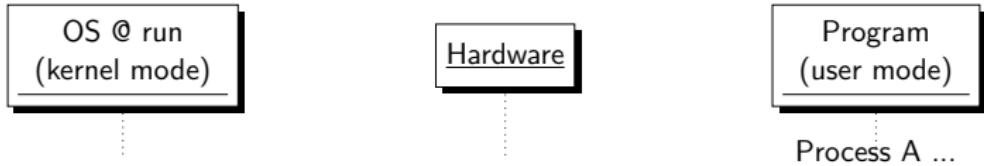
syscall handler

timer handler

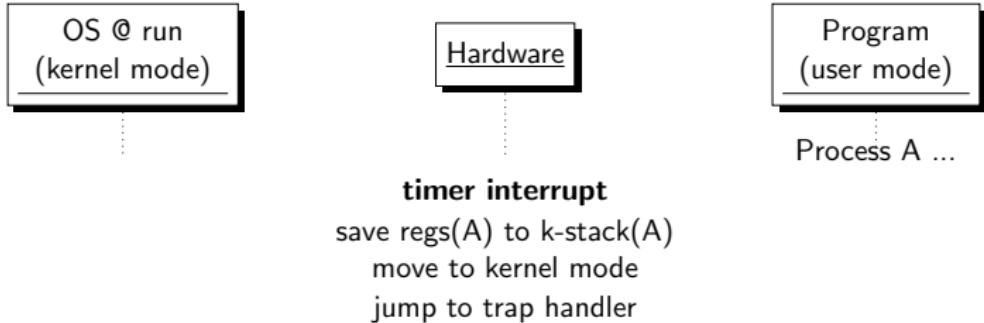
start interrupt timer

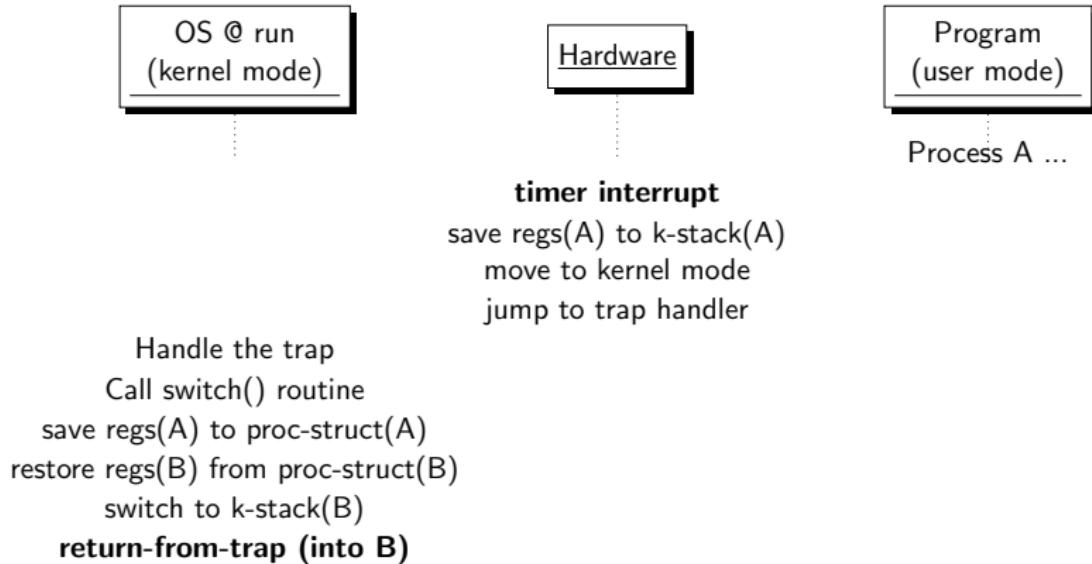
start timer

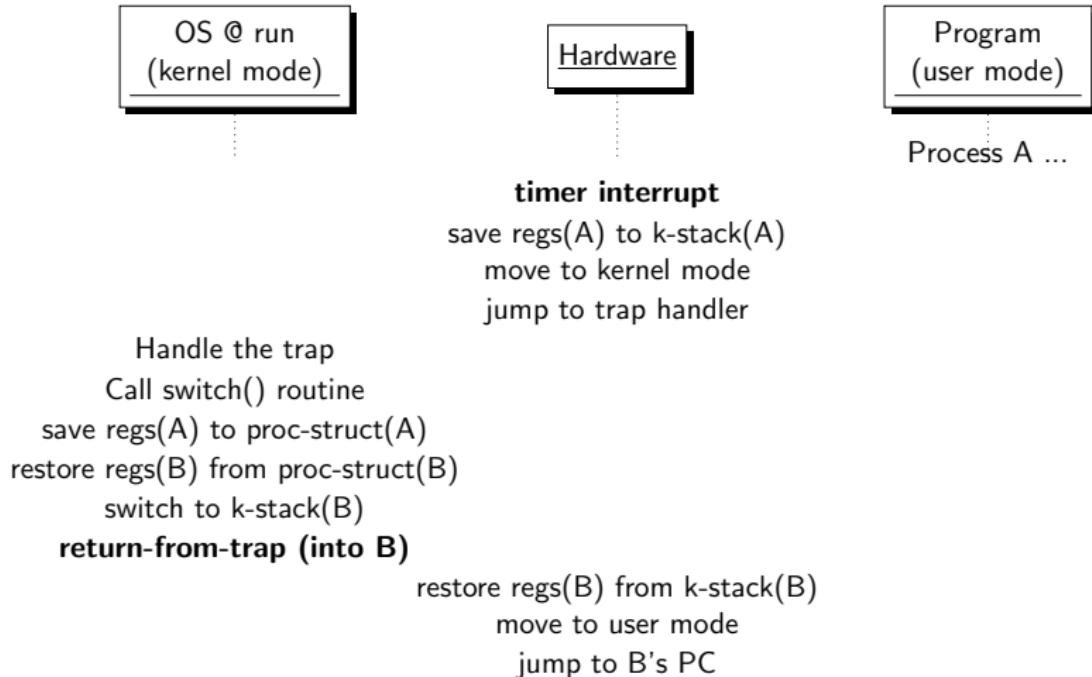
interrupt CPU in X ms

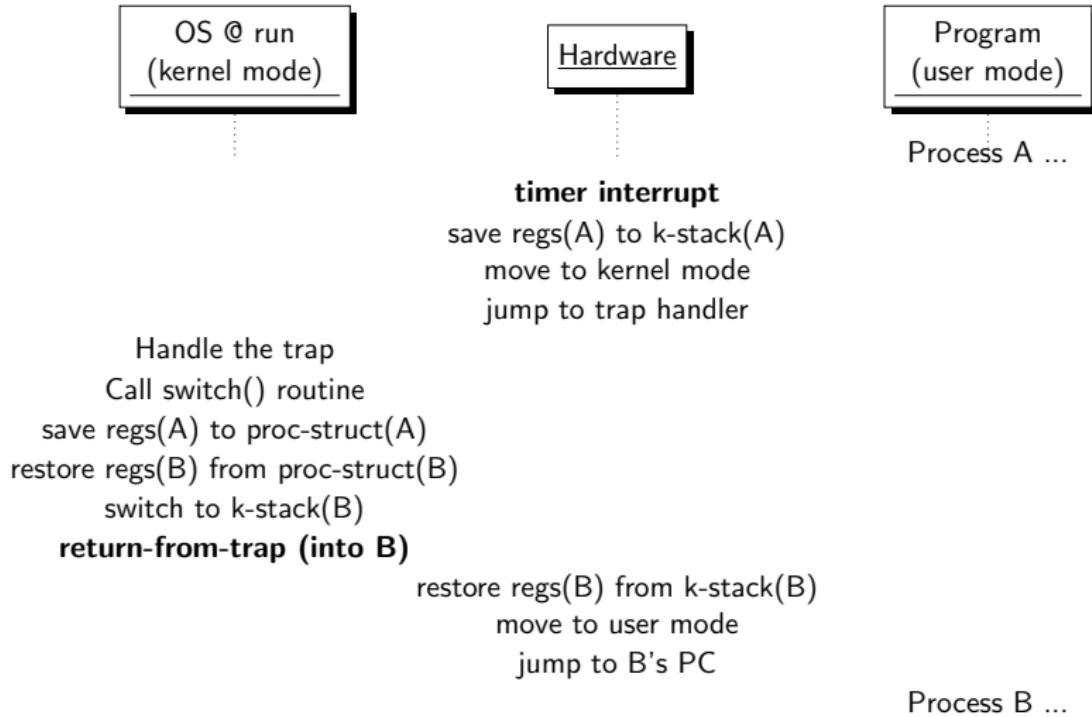


Process A ...









What happens if, during interrupt or trap handling, another interrupt occurs?

- ▶ Idea: Disable interrupts
- ▶ Locking mechanisms
- ▶ Will be discussed in lectures on concurrency

- The CPU should support at least two modes of execution: a restricted **user mode** and a privileged (non-restricted) **kernel mode**.
- Typical user applications run in user mode, and use a **system call** to **trap** into the kernel to request operating system services.
- The trap instruction saves register state carefully, changes the hardware status to kernel mode, and jumps into the OS to a pre-specified destination: the **trap table**.
- When the OS finishes servicing a system call, it returns to the user program via another special **return-from-trap** instruction, which reduces privilege and returns control to the instruction after the trap that jumped into the OS.
- The trap tables must be set up by the OS at boot time, and make sure that they cannot be readily modified by user programs. All of this is part of the **limited direct execution** protocol which runs programs efficiently but without loss of OS control.
- Once a program is running, the OS must use hardware mechanisms to ensure the user program does not run forever, namely the **timer interrupt**. This approach is a **non-cooperative** approach to CPU scheduling.
- Sometimes the OS, during a timer interrupt or system call, might wish to switch from running the current process to a different one, a low-level technique known as a **context switch**.