

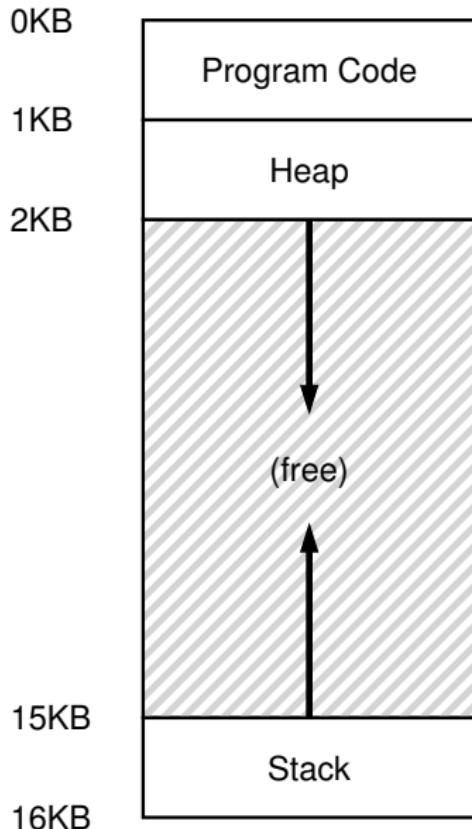
CSL 301

OPERATING SYSTEMS

Lecture 7

Address Space

Instructor
Dr. Dhiman Saha

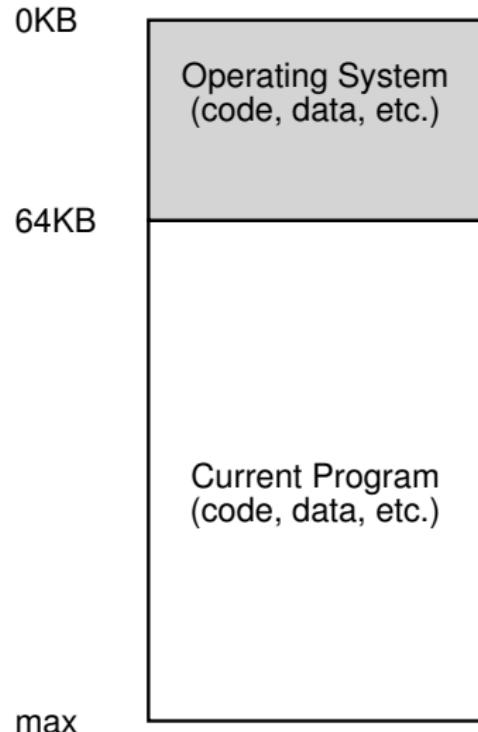


How to virtualize memory?

How can the OS build this abstraction of a private, potentially large address space for multiple running processes (all sharing memory) on top of a **single**, physical memory?

A Historical Perspective

- ▶ Early Systems
- ▶ The OS \equiv a set of routines
- ▶ Only one running program
- ▶ Few illusions!



- ▶ Multiple processes
- ▶ Ready to run at a given time
- ▶ The OS had a new task
- ▶ Switch between them

Goals

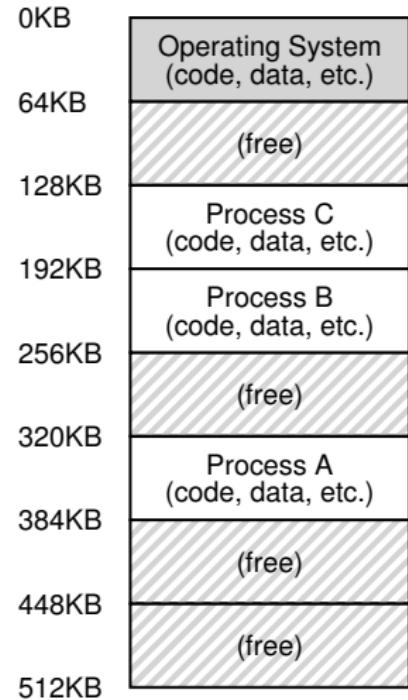
- ▶ Effective utilization
- ▶ Efficiency

- ▶ Interactivity
- ▶ Many users concurrently using a machine

Can you suggest one possible solution?

- ▶ How did we do this in case of CPU virtualization?
- ▶ Can we follow the same approach?
- ▶ What are the issues?

The physical memory allocated to a process could be non-contiguous too



Task

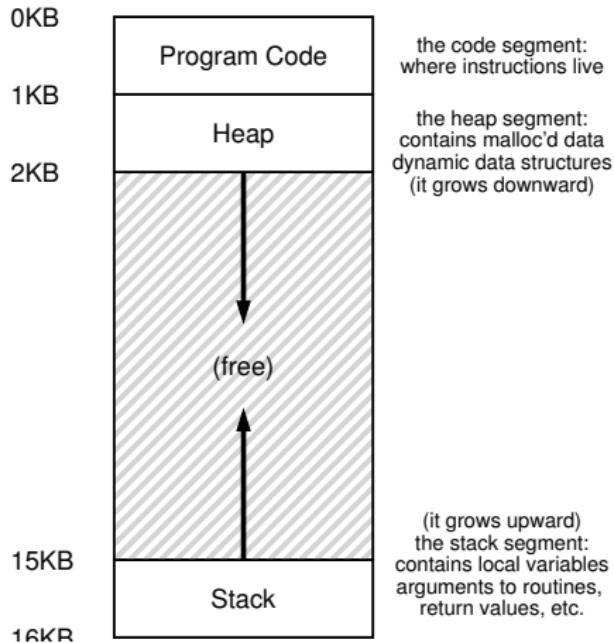
To create an easy to use abstraction of physical memory

Address Space

The running program's view of memory in the system.

- ▶ A process has a set of addresses that map to bytes
- ▶ This set is called address space
- ▶ How can we provide a private address space?

What is in the address space?

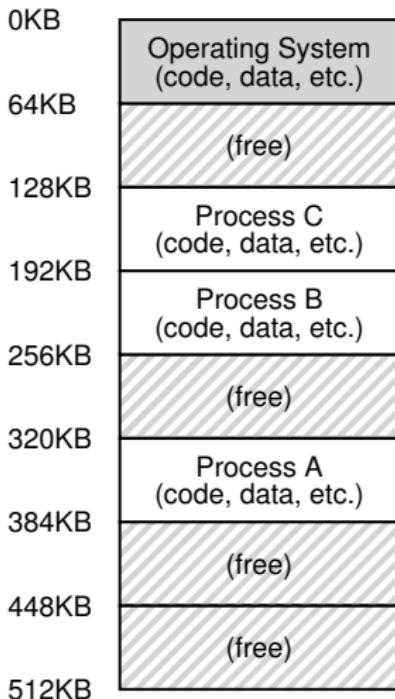


```
int x;  
int main(int argc, char* argv[]) {  
    int y;  
    int *z = malloc(sizeof(int));  
}
```

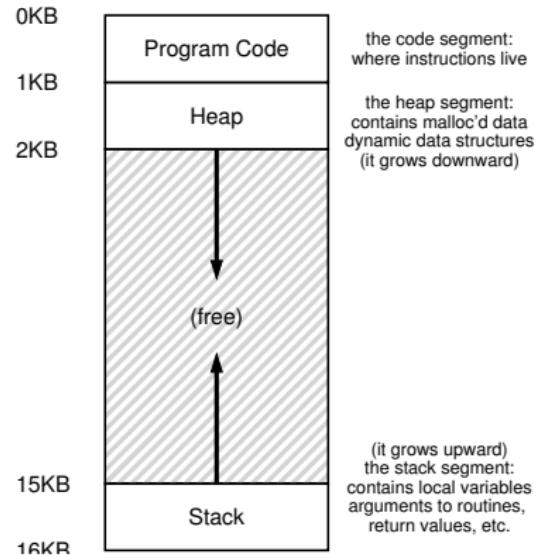
- ▶ x → code
- ▶ main → code
- ▶ y → stack
- ▶ z → heap

How are code, stack, heap arranged in this address space? Why?

Process A tries to perform a load at address 0



Actual View



View of Process A

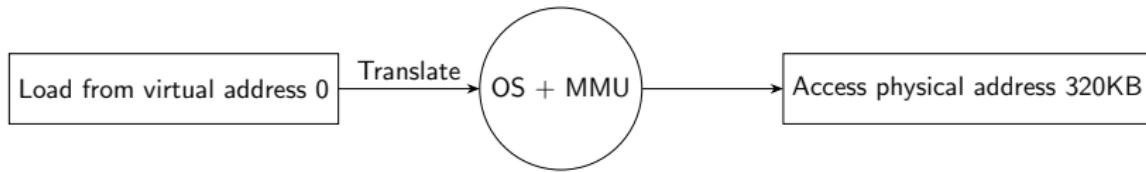
The Crux: How to Virtualize Memory?

The Core Problem

How can the OS build this private, clean-looking address space for multiple processes on top of a single, messy physical memory?

The Answer: Virtualization!

- ▶ The address space seen by the program is **virtual**.
- ▶ When a program tries to access an address (e.g., address '0'), the OS, with hardware help, translates this **virtual address** to a **physical address**.



- ▶ Virtual addresses (VA) to physical addresses (PA)
- ▶ CPU issues loads/stores to VA
- ▶ Memory hardware accesses PA
- ▶ OS allocates memory and tracks location of processes
- ▶ Memory Management Unit (MMU): memory hardware that does the translation
- ▶ OS makes the necessary information available

- ▶ One of the primary things that the OS achieves by abstracting out the physical memory
- ▶ What is the implication?

HW

Short Notes

- ▶ Microkernel
- ▶ Monolithic Kernel

Goals of a Virtual Memory System

► Transparency

- The process should be completely unaware that memory is being virtualized.
- It should behave as if it has its own private memory. The OS and hardware handle the magic behind the scenes.

Goals of a Virtual Memory System

► Transparency

- The process should be completely unaware that memory is being virtualized.
- It should behave as if it has its own private memory. The OS and hardware handle the magic behind the scenes.

► Efficiency

- Virtualization should be as fast as possible (time) and not consume too much memory for its own data structures (space).
- This relies heavily on hardware support (e.g., TLBs).

Goals of a Virtual Memory System

► Transparency

- ▶ The process should be completely unaware that memory is being virtualized.
- ▶ It should behave as if it has its own private memory. The OS and hardware handle the magic behind the scenes.

► Efficiency

- ▶ Virtualization should be as fast as possible (time) and not consume too much memory for its own data structures (space).
- ▶ This relies heavily on hardware support (e.g., TLBs).

► Protection

- ▶ Protect processes from each other. A bug in Process A should not corrupt Process B.
- ▶ Protect the OS from user processes.
- ▶ This provides the fundamental principle of **isolation**.

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    printf("location of code : %p\n", main);
    printf("location of heap : %p\n", malloc(1));
    int x = 3;
    printf("location of stack : %p\n", &x);
    return 0;
}
```

```
location of code : 0x55658716e6fa
location of heap : 0x556587aa2670
location of stack : 0x7ffdcd2c26b4
```

malloc() and free

Are these system calls?

- ▶ malloc library manages space within virtual address space
- ▶ Depends on some other system calls
- ▶ brk/sbrk
- ▶ Alternatively, mmap()

- ▶ Common mistakes while handling memory
- ▶ Using tools like valgrind and purify