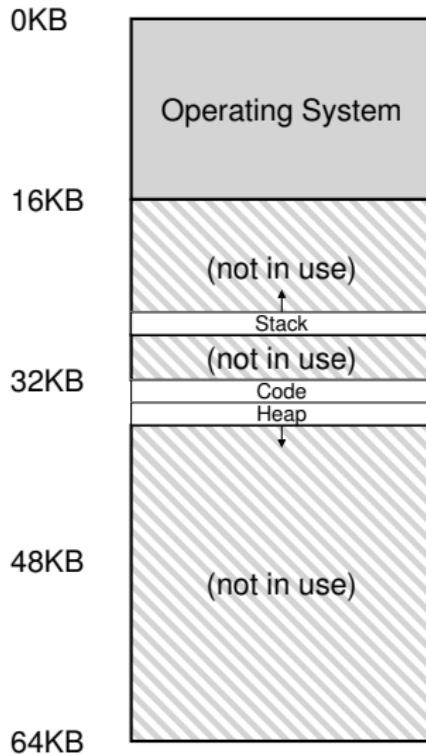


# CSL 301

## OPERATING SYSTEMS



Lecture 7  
Address Translation  
Base/Bounds  
Segmentation

Instructor  
Dr. Dhiman Saha

## The Crux: Virtualizing Memory

## The Core Challenge

How can we build an **efficient** and **flexible** virtualization of memory?

# The Crux: Virtualizing Memory

## The Core Challenge

How can we build an **efficient** and **flexible** virtualization of memory?

- ▶ **Efficiency:** We must rely on hardware support. We can't have the OS intervene on every memory access.
- ▶ **Control:** The OS must ensure processes cannot access memory outside their own address space (isolation).
- ▶ **Flexibility:** Programs should be able to use their address space in any way they like (e.g., have large, sparse data structures).

- 1 User's address space is placed contiguously in physical memory.
- 2 Size of the address space is not too big
- 3 It is less than the size of physical memory
- 4 Each address space is exactly the same size

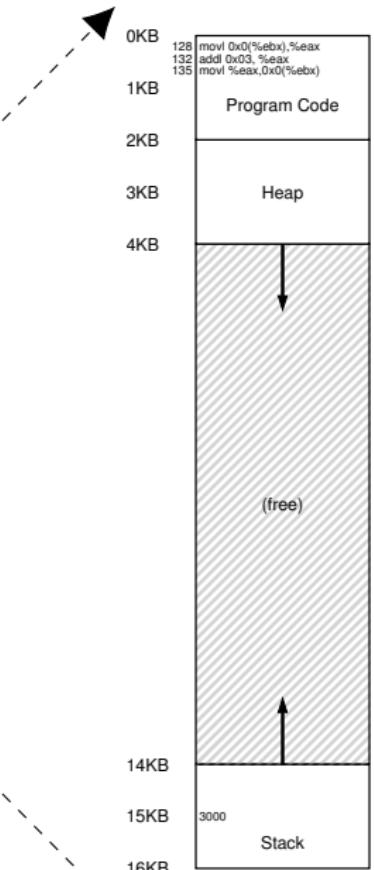
```
void func() {  
    int x = 3000;  
    x = x + 3;      // point of interest  
...  
}
```

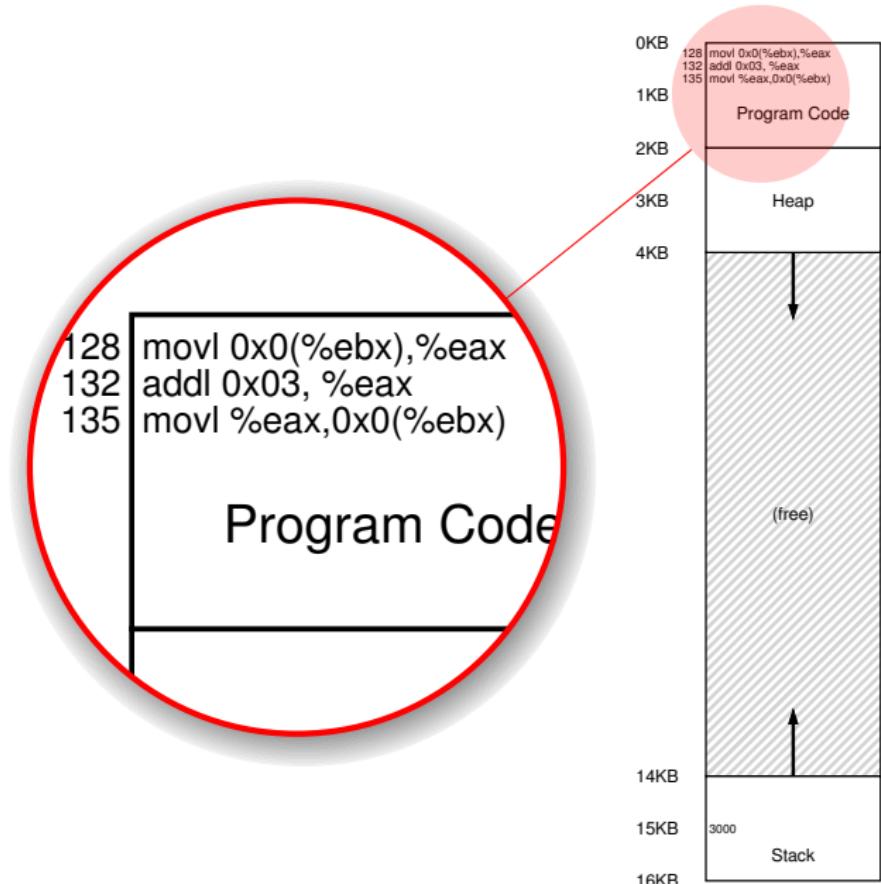
```
void func() {  
    int x = 3000;  
    x = x + 3;      // point of interest  
    ...  
}
```

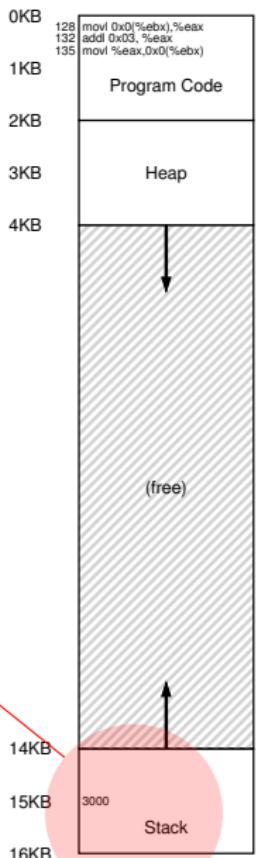
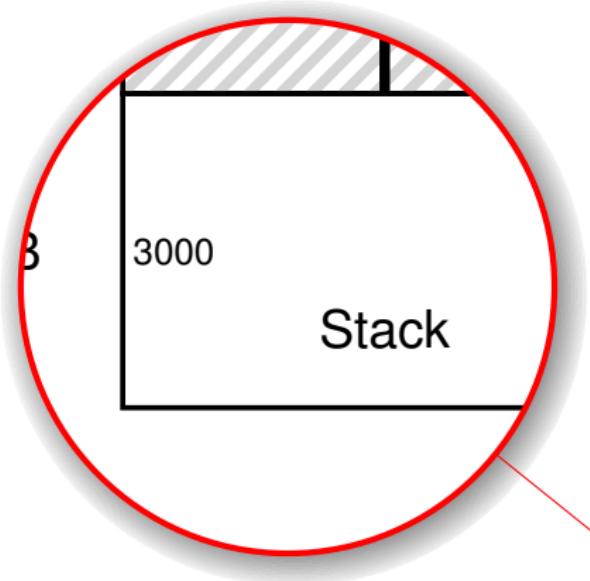
► Using objdump (in x86 assembly)

```
128: movl 0x0(%ebx), %eax      ;load 0+ebx into eax  
132: addl $0x03, %eax        ;add 3 to eax register  
135: movl %eax, 0x0(%ebx)     ;store eax back to mem
```

## The Address Space







- ▶ 128: `movl 0x0(%ebx), %eax ;load 0+ebx into eax`
- ▶ 132: `addl $0x03, %eax ;add 3 to eax register`
- ▶ 135: `movl %eax, 0x0(%ebx) ;store eax back to mem`

- ▶ Fetch instruction at address 128
- ▶ Execute this instruction (load from address 15 KB)
- ▶ Fetch instruction at address 132
- ▶ Execute this instruction (**no memory reference**)
- ▶ Fetch the instruction at address 135
- ▶ Execute this instruction (store to address 15 KB)

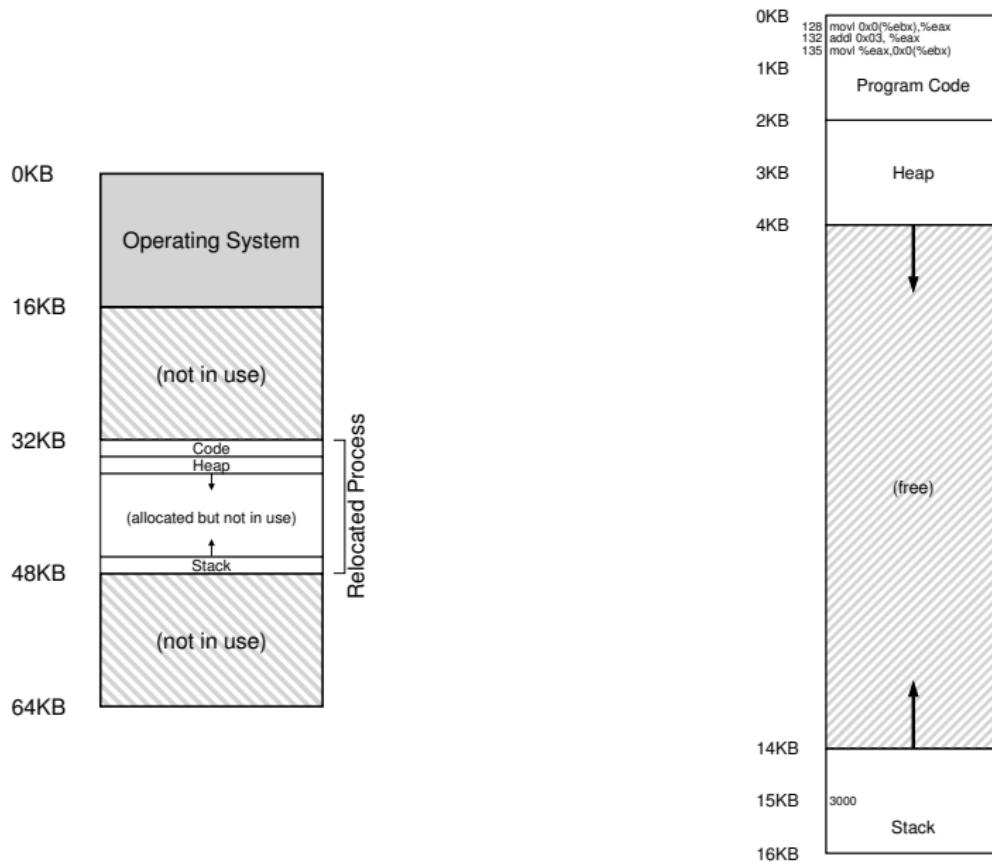
- ▶ 128: `movl 0x0(%ebx), %eax ;load 0+ebx into eax`
- ▶ 132: `addl $0x03, %eax ;add 3 to eax register`
- ▶ 135: `movl %eax, 0x0(%ebx) ;store eax back to mem`

- ▶ Fetch instruction at address 128
- ▶ Execute this instruction (load from address 15 KB)
- ▶ Fetch instruction at address 132
- ▶ Execute this instruction (**no memory reference**)
- ▶ Fetch the instruction at address 135
- ▶ Execute this instruction (store to address 15 KB)

- ▶ 128: `movl 0x0(%ebx), %eax ;load 0+ebx into eax`
- ▶ 132: `addl $0x03, %eax ;add 3 to eax register`
- ▶ 135: `movl %eax, 0x0(%ebx) ;store eax back to mem`

- ▶ Fetch instruction at address 128
- ▶ Execute this instruction (load from address 15 KB)
- ▶ Fetch instruction at address 132
- ▶ Execute this instruction (**no memory reference**)
- ▶ Fetch the instruction at address 135
- ▶ Execute this instruction (store to address 15 KB)

# Physical Memory with a Single Relocated Process



## Attempt #1: Software-Based Relocation

Static Relocation

## Idea

Rewrite the program itself before **loading** it as a process

- ▶ Using software support: **loader**
- ▶ The **loader** takes an executable that is about to be run
- ▶ **Rewrites its addresses**
- ▶ To the desired **offset** in physical memory

## Classwork

- ▶ How would it effect our example program?
- ▶ What if there are multiple processes?

Why?

- ▶ Protection
- ▶ Re-Relocation

# Attempt #2: Hardware-Based Relocation The Base Register

Dynamic Relocation

## Idea

- ▶ Address translation by adding a fixed offset.
  - ▶ Offset stored in *Base* Register
  - ▶ Base register has different value for each process
- 
- ▶ OS tells the hardware the base (starting address)
  - ▶ Memory hardware calculates PA from VA
  - ▶ “dynamic relocation”

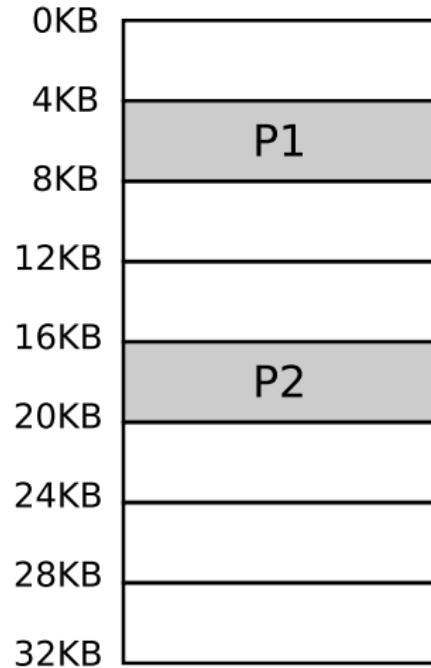
$$\text{physical address} = \text{virtual address} + \text{base}$$

## Note

Each program is written and compiled as if it is loaded at address zero.

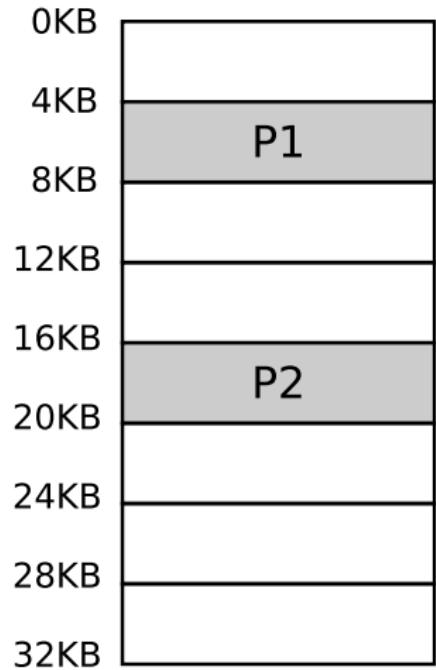
## Let us try an example

- ▶ Two Process Scenario
- ▶ To Do: Address Translation



How?

- ▶ Protection



# Attempt #3: Hardware-Based Relocation Base + Bounds

Dynamic Relocation

## Idea

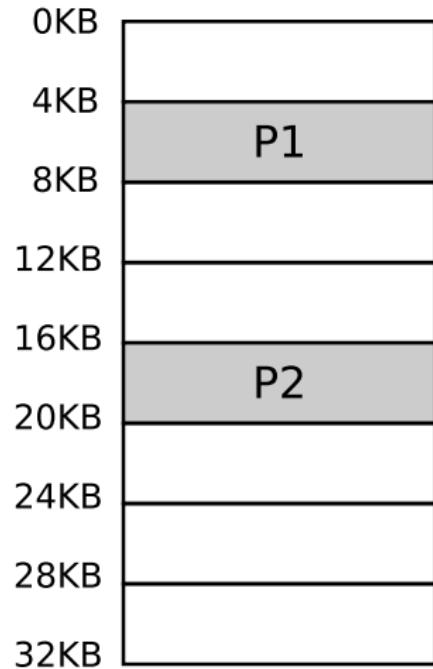
- ▶ “Bound” the address space
  - ▶ The largest addressable physical address for a process
  - ▶ Stored in *Bounds (limit)* register
- 
- ▶ Base → translate the address
  - ▶ Bounds → ensure physical address lies within address space

## Classwork

- ▶ What can bounds-register actually store?
- ▶ Where are these registers located?

## Let us try an example

- ▶ Two Process Scenario
- ▶ To Do: Illegal Memory Access



- ▶ A special data-structure used by OS
- ▶ To track which parts of free memory are not in use
- ▶ Simply is a list of the ranges of the physical memory which are not currently in use

Hardware Requirements	Notes
Privileged mode	<i>Needed to prevent user-mode processes from executing privileged operations</i>
Base/bounds registers	<i>Need pair of registers per CPU to support address translation and bounds checks</i>
Ability to translate virtual addresses and check if within bounds	<i>Circuitry to do translations and check limits; in this case, quite simple</i>
Privileged instruction(s) to update base/bounds	<i>OS must be able to set these values before letting a user program run</i>
Privileged instruction(s) to register exception handlers	<i>OS must be able to tell hardware what code to run if exception occurs</i>
Ability to raise exceptions	<i>When processes try to access privileged instructions or out-of-bounds memory</i>

OS Requirements	Notes
Memory management	<i>Need to allocate memory for new processes; Reclaim memory from terminated processes; Generally manage memory via free list</i>
Base/bounds management	<i>Must set base/bounds properly upon context switch</i>
Exception handling	<i>Code to run when exceptions arise; likely action is to terminate offending process</i>

## Classwork

Can the OS perform **address space relocation** when a process not running? How?

- ▶ Look at LDE protocol with dynamic relocation in OSTEP book.

# Base & Bounds: Translation and Protection

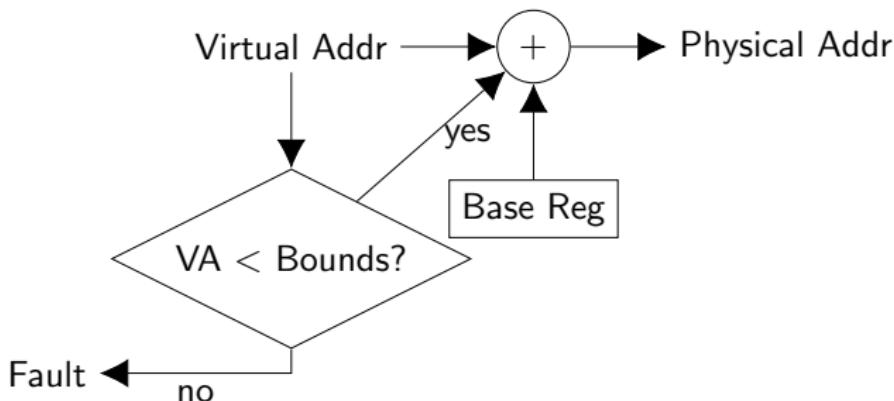
For every memory access, the hardware performs two steps:

- ▶ **Protection Check:** Is the access legal?

```
if (virtual_addr >= bounds) -> raise exception;
```

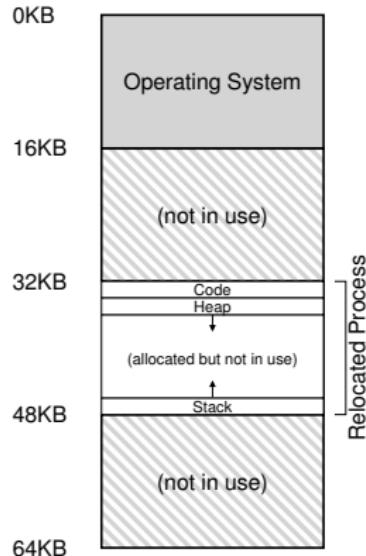
- ▶ **Address Translation:** If legal, calculate the physical address.

```
physical_addr = virtual_addr + base;
```



# Issues?

# Base + Bounds

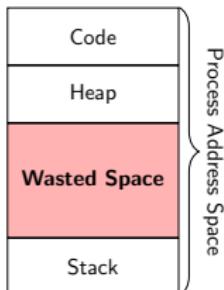


- ▶ 1. Internal Fragmentation
- ▶ 2. What if address space does not fit into memory?

# The Problem: Internal Fragmentation

The base-and-bounds approach is simple, but wasteful.

- ▶ The entire address space, from address 0 to the top of the stack, must be mapped to a *contiguous* chunk of physical memory.
- ▶ The large, unused space between the heap and the stack is also allocated.
- ▶ This waste is called **internal fragmentation**.



This makes it hard to run programs if their full address space doesn't fit in memory.

# Attempt #4: Hardware-Based Relocation Segmentation

Generalized Base/Bounds

## Generalized Base/Bounds

Idea

Instead of having just one base and bounds pair in our MMU, why not have a base and bounds pair per logical **segment** of the address space?

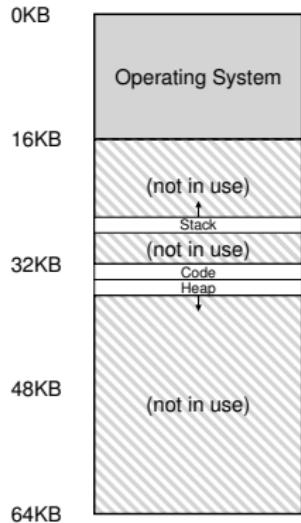
## Segment

A segment is just a **contiguous** portion of the address space of a particular length.

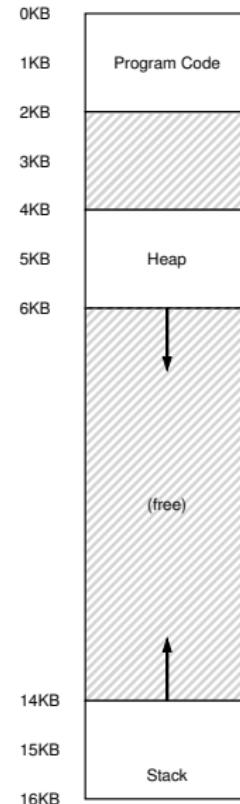
- ▶ Code
- ▶ Stack
- ▶ Heap

Three logically-different segments. How to utilize this setting?

# Placing Segments In Physical Memory



Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K

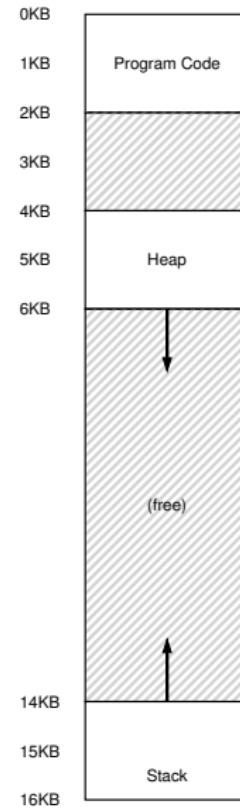
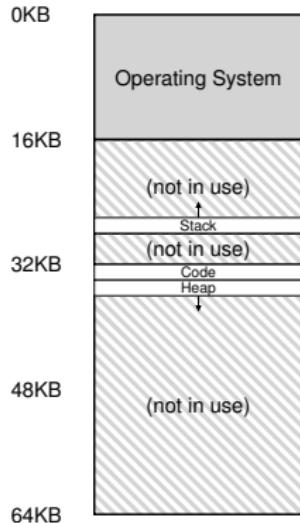


Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K

- ▶ A set of three base and bounds register pairs

# Address Translation

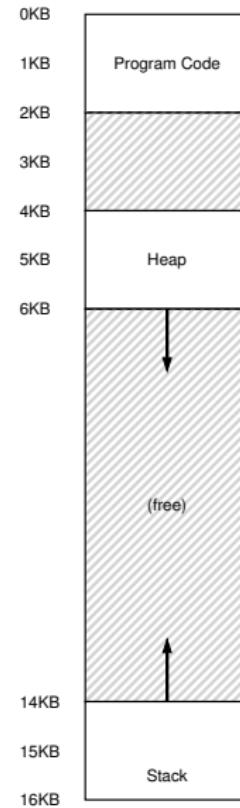
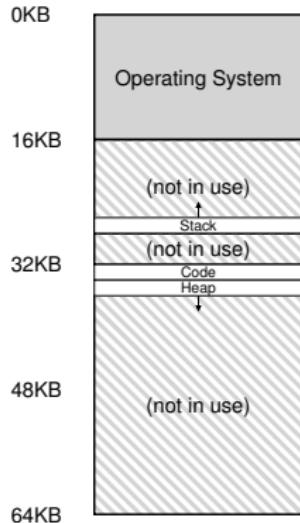
Let us try an example



► VA: 135 PA: \_\_\_\_\_ ?

# Address Translation

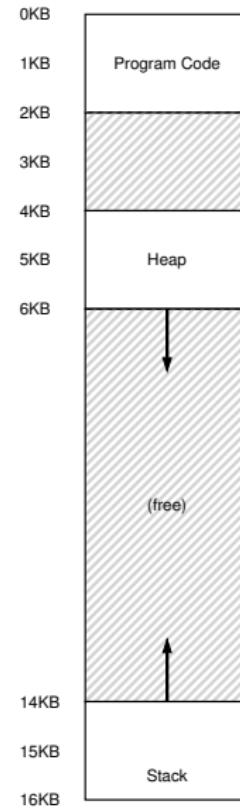
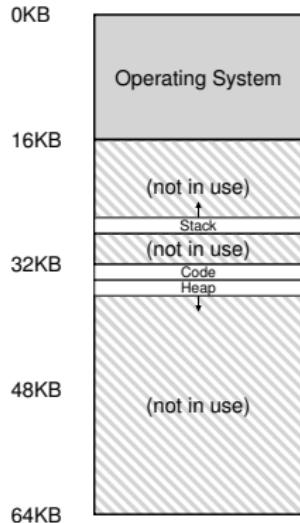
Let us try an example



- ▶ VA: 4400 PA: \_\_\_\_\_ ?

# Address Translation

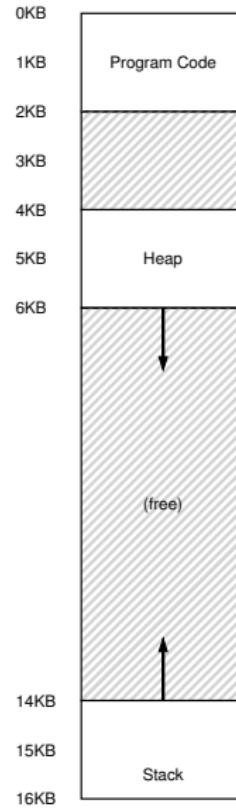
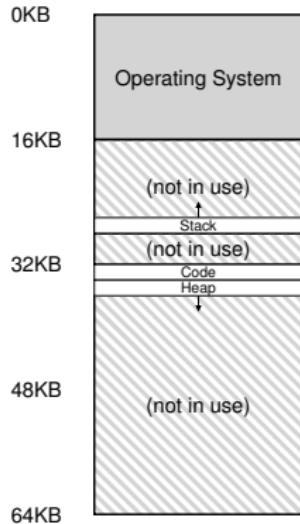
Let us try an example



- VA: 7KB PA: \_\_\_\_\_ ?

# Address Translation

Let us try an example

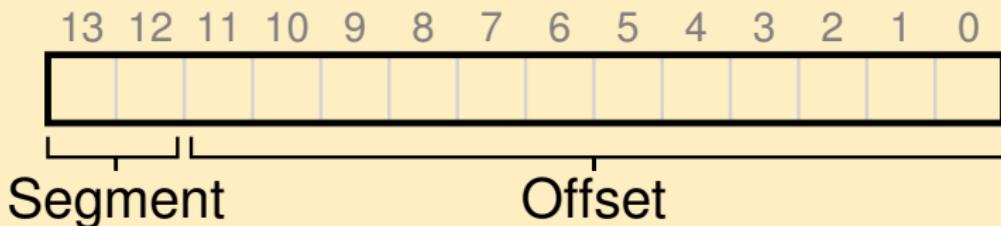


- ▶ VA: 7KB PA: \_\_\_\_\_ ?
- ▶ Segmentation Fault

# Which Segment Are We Referring To?

# Which Segment Are We Referring To?

Explicit Approach



Classwork

Calculate this for VA:4400

```
// get top 2 bits of 14-bit VA
Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT
// now get offset
Offset = VirtualAddress & OFFSET_MASK
if (Offset >= Bounds[Segment])
    RaiseException(PROTECTION_FAULT)
else
    PhysAddr = Base[Segment] + Offset
    Register = AccessMemory(PhysAddr)
```

## Classwork

### Calculate

- ▶ SEG\_MASK
- ▶ SEG\_SHIFT
- ▶ OFFSET\_MASK

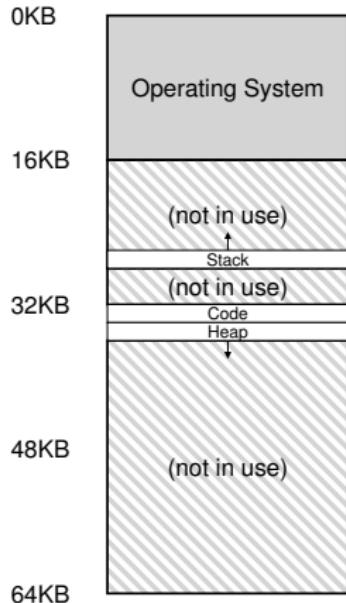
# Which Segment Are We Referring To?

## Implicit Approach

The hardware determines the segment by noticing **how the address was formed**

- ▶ Address generated from PC  $\implies$  code segment
- ▶ Address based of stack pointer  $\implies$  stack segment
- ▶ Otherwise  $\implies$  heap segment

# Did we forget the stack?



Recall  
Stack grows backwards

Negative-Growth Support

Segment	Base	Size	Grows Positive?
Code	32K	2K	1
Heap	34K	2K	1
Stack	28K	2K	0

Segment Registers

HomeWork

How will address translation take place now?