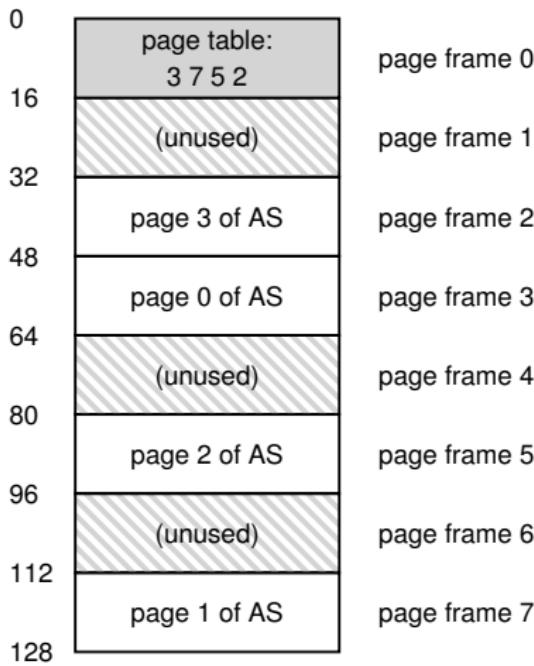


# CSL 301

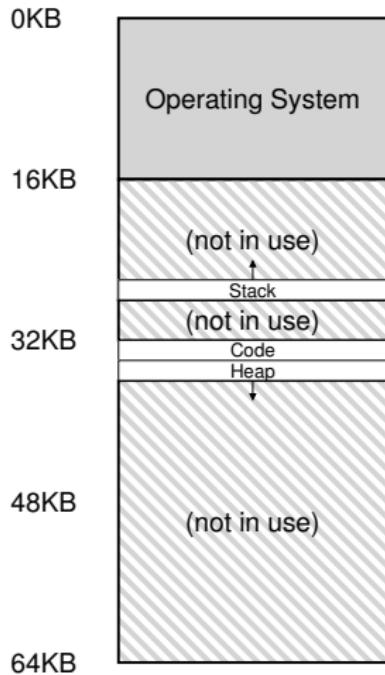
## OPERATING SYSTEMS



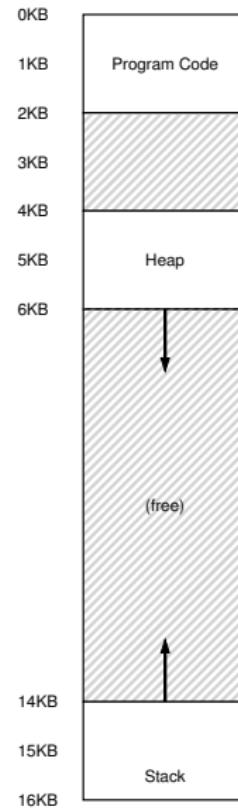
Lecture 10  
Segmentation (Contd.) +  
Paging

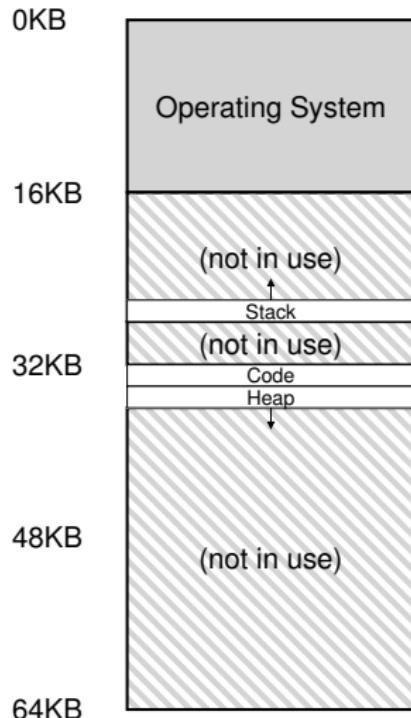
Instructor  
Dr. Dhiman Saha

# Recall



# Segmentation





- ▶ Pieces of the address space are relocated into physical memory as the system runs
- ▶ Implies huge savings of physical memory w.r.t single base/bound register approach
- ▶ Note **unused space** saved between heap and stack  
**need no longer be allocated**

## Idea

Share certain memory segments between address spaces

- ▶ For instance **code sharing**
- ▶ Hardware support: **protection bits**

Segment	Base	Size	Grows Positive?	Protection
Code	32K	2K	1	Read-Execute
Heap	34K	2K	1	Read-Write
Stack	28K	2K	0	Read-Write

- ▶ New task: Hardware has to check if a particular access is permissible

## Fine/Coarse Grained

## Granularity of segmentation

## Coarse-grained

- ▶ Code/Stack/Heap
  - ▶ Relatively large segments
  - ▶ Hence coarse-grained
- ▶ Idea of **segment table** (stored in memory)

## Fine-grained

- ▶ Idea of smaller segments
- ▶ Many segments  $\implies$  more H/W support

## How does this help?

With fine-grained segments, the OS could better learn about which segments are in use and which are not.

What should the OS do on a context switch?

- ▶ The segment registers must be saved and restored

Second issue is managing free space in physical memory

- ▶ The OS has to be able to find space in physical memory for the segments of a new process

### Note

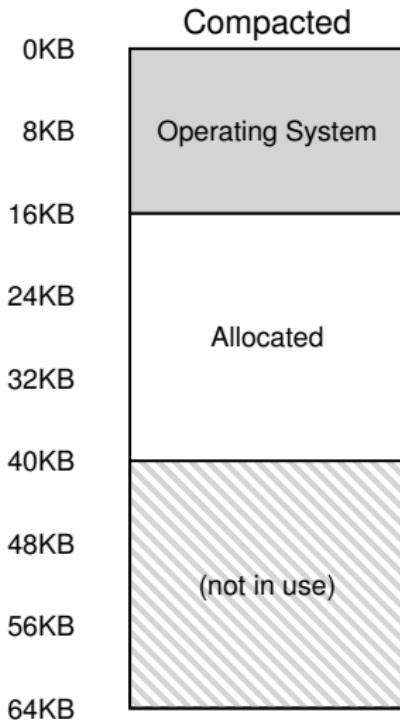
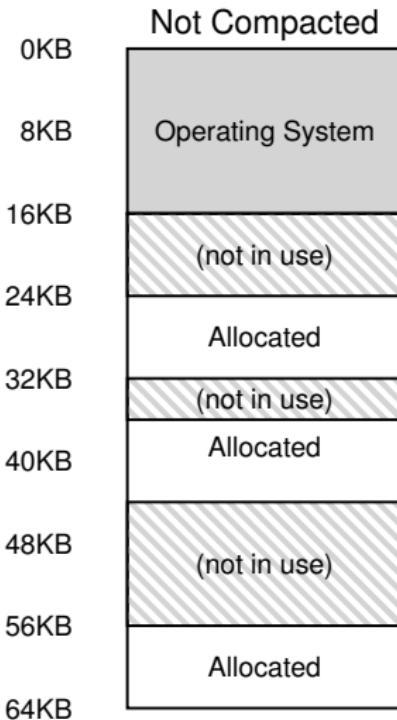
We have a number of segments per process, and each segment might be a different size.

### External Fragmentation

Physical memory quickly becomes **full of little holes of free space**, making it difficult to allocate new segments, or to grow existing ones.

## Solution #1: Compaction

## External Fragmentation



The OS cannot (left)/ can (right) satisfy a 20KB request.

### Note

- ▶ Compaction is memory-intensive
- ▶ And what about the CPU activity during compaction?

### Free-list management algorithm

- ▶ Best-fit
- ▶ Worst-fit
- ▶ Buddy Algorithm (covered later in the course)

Can you get rid of external fragmentation?

- ▶ External fragmentation
- ▶ Not flexible enough to support our fully generalized, sparse address space

## Example

If we have a large but sparsely-used heap all in one logical segment, the entire heap must still reside in memory in order to be accessed.

- ▶ Actual address space usage is not taken into account

- 1 User's address space is placed contiguously in physical memory.
- 2 Size of the address space is not too big
- 3 It is less than the size of physical memory
- 4 Each address space is exactly the same size

# Attempt #5: Hardware-Based Relocation Paging

## Idea

Divide address space into fixed-sized units, called **pages**.

- ▶ Same for physical memory in terms of **page-frames**

$$\text{page}_{\text{virtual-memory}} \equiv \text{page-frame}_{\text{physical-memory}}$$

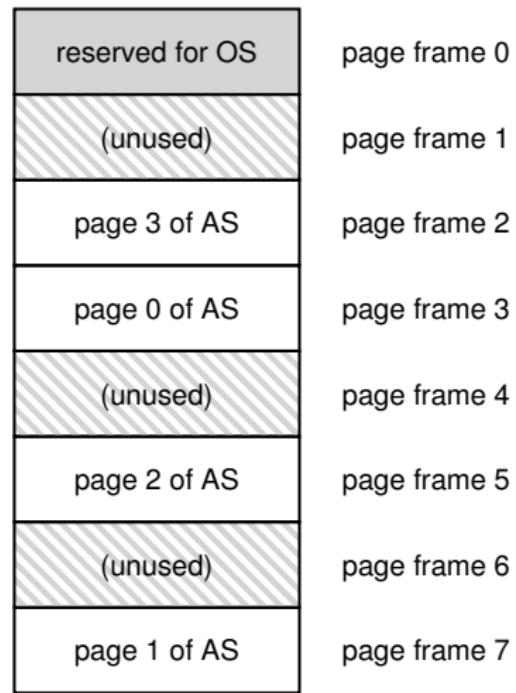
# Pages



# Page Frames



A 64-byte Address Space



A 128-Byte Physical Memory

# Memory Virtualization with Paging

- ▶ How can we virtualize memory with pages, so as to avoid the problems of segmentation?
- ▶ What are the basic techniques?
- ▶ How do we make those techniques work well, with minimal space and time overheads?

## Flexibility

- ▶ Better abstraction of the address space
- ▶ Does not rely on **how** a process uses the address space
- ▶ Recall heap, stack **growth** in segmentation
- ▶ **Simplicity** of free-space management

## How?

No external fragmentation

## Task

To record where each **virtual page** of the **address space** is placed in physical memory

## Solution

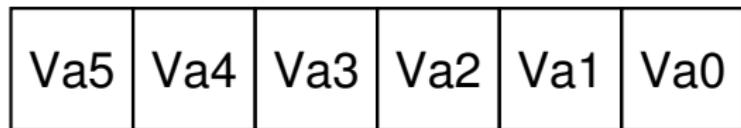
A **per-process** data structure known as a **page table**

- ▶ Basically stores the address translations for each of the virtual pages of the address space
- ▶ Note that this a **per-process** feature
- ▶ What changes during a context-switch?

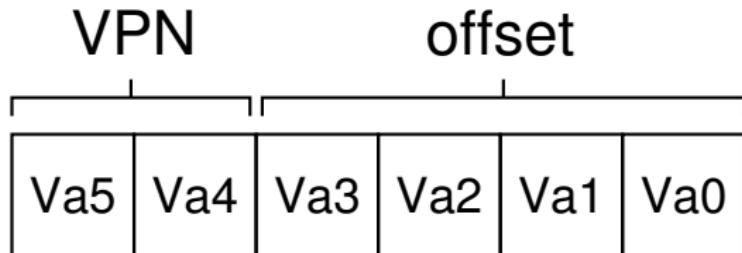
- ▶ `movl <virtual address>, %eax`
- ▶ Recal: 64 byte address space

Va5	Va4	Va3	Va2	Va1	Va0
-----	-----	-----	-----	-----	-----

- ▶ `movl <virtual address>, %eax`
- ▶ Recal: 64 byte address space

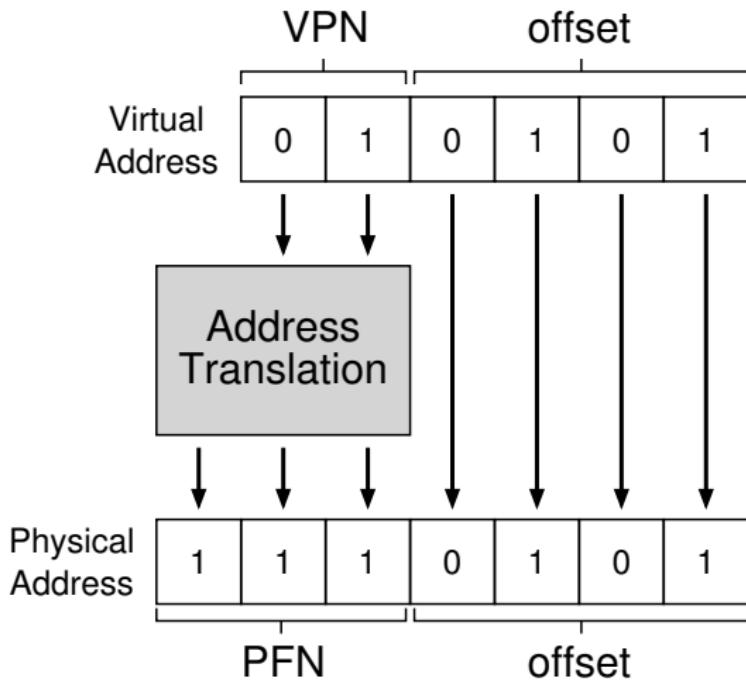


- ▶ Virtual Address = VPN + Offset



# Address-Translation (Classwork)

movl 21, %eax



- ▶ Where are these page tables stored?
- ▶ What are the typical contents of the page table?
- ▶ How big are the tables?
- ▶ Does paging make the system (too) slow?

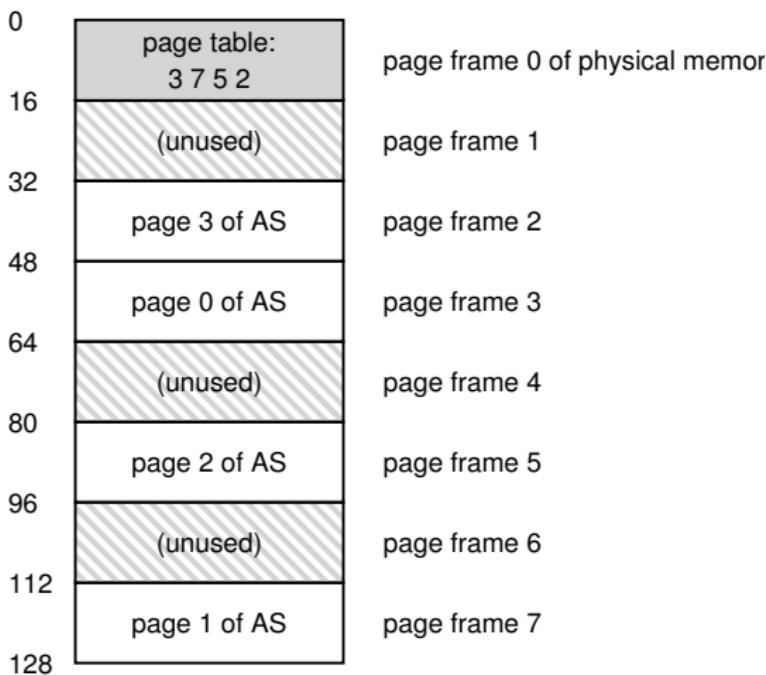
# Where are page tables stored?

Should we use hardware support like we have done so far?

Let us do the math:

- ▶ Address Space: 32-bit
- ▶ Page Size (say): 4KB
- ▶ VPN Size: \_\_\_\_\_
- ▶ Total number of translations: \_\_\_\_\_
- ▶ Typical **Page Table Entry** Size: 4 bytes
- ▶ Page Table Size: \_\_\_\_\_
- ▶ Multiply by #processes. **So where should we store PT?**

# Where are page tables stored?



Page Table in Kernel Physical Memory

# What's Actually In The Page Table?

## Page Table Organization

What is the simplest data structure that comes to mind?

PTE Contents:

- ▶ Valid bit
- ▶ Protection bits
- ▶ Present bit
- ▶ Dirty bit
- ▶ Reference bit (a.k.a. accessed bit) and some other bits
- ▶ PFN

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
PFN														G	PAT	D	A	PCD	PWT	U/S	R/W	P									

An x86 Page Table Entry (PTE)

# Does paging make the system (too) slow?

Let us see what happens?

Our Example

```
movl 21, %eax
```

- ▶ Step 1: Where is the page table of the process in memory?

Page-table base register

Contains the physical address of the starting location of the page table

- ▶ Step 2: To get correct PTE for translation hardware will do:

```
VPN    = (VirtualAddress & VPN_MASK) >> SHIFT
```

```
PTEAddr = PageTableBaseReg + (VPN * sizeof(PTE))
```

# Does paging make the system (too) slow?

- ▶ Step 3: The hardware will fetch the PTE from memory, extract the PFN

## Note

This is a memory access that the system cannot avoid!!!

- ▶ Step 4: The hardware will do the final translation:

```
offset    = VirtualAddress & OFFSET_MASK  
PhysAddr  = (PFN << SHIFT) | offset
```

- ▶ Step 5: The hardware can fetch the desired data from memory and put it into register eax.

```
1 // Extract the VPN from the virtual address
2 VPN = (VirtualAddress & VPN_MASK) >> SHIFT
3
4 // Form the address of the page-table entry (PT
5 PTEAddr = PTBR + (VPN * sizeof(PTE))
6
7 // Fetch the PTE
8 PTE = AccessMemory(PTEAddr)
9
10 // Check if process can access the page
11 if (PTE.Valid == False)
12     RaiseException(SEGMENTATION_FAULT)
13 else if (CanAccess(PTE.ProtectBits) == False)
14     RaiseException(PROTECTION_FAULT)
15 else
16     // Access is OK: form physical address and
17     offset = VirtualAddress & OFFSET_MASK
18     PhysAddr = (PTE.PFN << PFN_SHIFT) | offset
19     Register = AccessMemory(PhysAddr)
```

## Space overheads

Issue 1

Storing PT in memory wastes valuable memory space.

## > Factor 2 slow down

Issue 2

For every memory reference, paging requires us to perform one extra memory reference in order to first fetch the translation from the page table

## Attempt #6: Hardware-Based Relocation Translation-lookaside Buffer

Faster Paging