

Pandas is a Python library used for working with data sets. It has functions for analyzing, cleaning, exploring, and manipulating data.

Installation of Pandas

```
!pip install pandas

Requirement already satisfied: pandas in
/usr/local/lib/python3.10/dist-packages (2.2.2)
Requirement already satisfied: numpy>=1.22.4 in
/usr/local/lib/python3.10/dist-packages (from pandas) (1.26.4)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.10/dist-packages (from pandas) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in
/usr/local/lib/python3.10/dist-packages (from pandas) (2024.2)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2-
>pandas) (1.17.0)
```

Import Pandas

```
import pandas as pd #importing pandas and providing it with an alias
```

Series

A **Pandas Series** is like a column in a table.

It is a *one-dimensional array* holding data of any type.

```
a = [1, 7, 2, 4, 9, 8]
```

```
myNum = pd.Series(a)
```

```
print(myNum)
```

```
0    1
1    7
2    2
3    4
4    9
5    8
```

```
dtype: int64
```

```
print(myNum[0])
```

```
print(myNum[5])
```

```
1
8
```

Labels

With the *index* argument, you can name your own labels.

```
a = [10, 17, 21]
myNum = pd.Series(a, index = ["a", "b", "c"])
print(myNum)
a      10
b      17
c      21
dtype: int64
print(myNum["a"])
10
```

Key/Value Objects as Series

```
running = {"day1": 2, "day2": 3, "day3": 5}
myRun = pd.Series(running)
print(myRun)
day1    2
day2    3
day3    5
dtype: int64
```

DataFrames

Data sets in Pandas are usually *multi-dimensional* tables, called **DataFrames**.

Series is like a column, a DataFrame is the whole table.

```
data = {
    "kilometers": [4, 3, 5],
    "duration": [50, 40, 45]
}
myRun = pd.DataFrame(data)
print(myRun)
   kilometers  duration
0            4         50
1            3         40
2            5         45
```

```
print(myRun.info())      # Overview of the dataset
print(myRun.describe())  # Summary statistics
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3 entries, 0 to 2
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   kilometers  3 non-null      int64
1   duration    3 non-null      int64
dtypes: int64(2)
memory usage: 176.0 bytes
None
```

	kilometers	duration
count	3.0	3.0
mean	4.0	45.0
std	1.0	5.0
min	3.0	40.0
25%	3.5	42.5
50%	4.0	45.0
75%	4.5	47.5
max	5.0	50.0

```
data = {'Name': ['Alice', 'Bob', 'Charlie'],
        'Age': [25, 30, 35]}
df = pd.DataFrame(data)
print(df)
```

	Name	Age
0	Alice	25
1	Bob	30
2	Charlie	35

Loading Data from a File

```
mydf =
pd.read_csv('https://raw.githubusercontent.com/gagan-iitb/DataAnalyticsAndVisualization/refs/heads/main/Lab-W25/dataset/names.csv')
```

Download CSV - [names.csv](#)

```
print(mydf.head())  # Display the first 5 rows
```

	Name	Age
0	Alice	25
1	Bob	30
2	Charlie	35
3	James	23
4	John	26

```
print(mydf.head(7)) # Display the first 7 rows
```

	Name	Age
0	Alice	25
1	Bob	30
2	Charlie	35
3	James	23
4	John	26
5	William	28
6	Caleb	25

```
print(mydf['Name']) # Single column
```

0	Alice
1	Bob
2	Charlie
3	James
4	John
5	William
6	Caleb
7	Helen

Name: Name, dtype: object

```
print(mydf[['Age', 'Name']]) # Multiple columns
```

	Age	Name
0	25	Alice
1	30	Bob
2	35	Charlie
3	23	James
4	26	John
5	28	William
6	25	Caleb
7	30	Helen

Filtering Rows

```
print(mydf[mydf['Age'] > 25])
```

	Name	Age
1	Bob	30
2	Charlie	35
4	John	26
5	William	28
7	Helen	30

Adding/Updating Columns

```
mydf['Salary'] = [50000, 60000, 50000, 50000, 30000, 70000, 90000, 80000]
print(mydf)
```

	Name	Age	Salary
0	Alice	25	50000
1	Bob	30	60000
2	Charlie	35	50000
3	James	23	50000
4	John	26	30000
5	William	28	70000
6	Caleb	25	90000
7	Helen	30	80000

Saving to a File

```
mydf.to_csv('myDataframe.csv', index=False)
```

Dropping Columns

```
mydf = mydf.drop('Salary', axis=1) # Drop column
print(mydf)
```

	Name	Age
0	Alice	25
1	Bob	30
2	Charlie	35
3	James	23
4	John	26
5	William	28
6	Caleb	25
7	Helen	30

#Create/Append two new columns named Marks, Department in mydf and display it

```
import pandas as pd
```

```
mydf['Marks'] = [85, 90, 88, 75, 80, 95, 85, 90]
mydf['Department'] = ['Design', 'Marketing', 'Development', 'Management', 'Marketing', 'Design', 'Development', 'Design']
print(mydf)
```

	Name	Age	Marks	Department
0	Alice	25	85	Design
1	Bob	30	90	Marketing
2	Charlie	35	88	Development
3	James	23	75	Management
4	John	26	80	Marketing
5	William	28	95	Design

6	Caleb	25	85	Development
7	Helen	30	90	Design

```
#Save the newly create mydf to a csv file. (Name of file =
myDataframe_YourIDNumber.csv)
mydf.to_csv('myDataframe_12340220.csv', index=False)
print("DataFrame has been saved to 'myDataframe_12340220.csv'.")
```

DataFrame has been saved to 'myDataframe_12340220.csv'.

```
#Filter all the rows where Age falls between 25-30.
filtered_df = mydf[(mydf['Age'] >= 25) & (mydf['Age'] <= 30)]
print(filtered_df)
```

	Name	Age	Marks	Department
0	Alice	25	85	Design
1	Bob	30	90	Marketing
4	John	26	80	Marketing
5	William	28	95	Design
6	Caleb	25	85	Development
7	Helen	30	90	Design

Unique() function

```
mydf.Age.unique()
array([25, 30, 35, 23, 26, 28])
```

Sorting

```
mydf.sort_values(by=['Age'])

{"summary":{"\n  \"name\": \"mydf\",\n  \"rows\": 8,\n  \"fields\": [\n    {\n      \"column\": \"Name\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 8,\n        \"samples\": [\n          \"Alice\",\n          \"Bob\",\n          \"James\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Age\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 3,\n        \"min\": 23,\n        \"max\": 35,\n        \"num_unique_values\": 6,\n        \"samples\": [\n          23,\n          25,\n          35\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Marks\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 6,\n        \"min\": 75,\n        \"max\": 95,\n        \"num_unique_values\": 6,\n        \"samples\": [\n          75,\n          85,\n          88\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Department\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 4,\n        \"samples\":
```

```

[\\n          \\\"Design\\\",\\n          \\\"Marketing\\\",\\n
\\\"Management\\\"\\n          ],\\n          \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n          }\\n          }\\n          ]\\n}\\",\"type\":\"dataframe\"}

#Sort mydf dataframe on the basis of Name,Marks.
mydf.sort_values(by=['Name', 'Marks'])

{\"summary\":{\"\\n  \\\"name\\\": \\\"mydf\\\",\\n  \\\"rows\\\": 8,\\n  \\\"fields\\\": [\\n
  {\\n    \\\"column\\\": \\\"Name\\\",\\n    \\\"properties\\\": {\\n
    \\\"dtype\\\": \\\"string\\\",\\n    \\\"num_unique_values\\\": 8,\\n
    \\\"samples\\\": [\\n      \\\"Bob\\\",\\n      \\\"James\\\",\\n
    \\\"Alice\\\"\\n    ],\\n    \\\"semantic_type\\\": \\\"\\\",\\n
    \\\"description\\\": \\\"\\\"\\n    }\\n    },\\n    {\\n      \\\"column\\\":
    \\\"Age\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n
    \\\"std\\\": 3,\\n        \\\"min\\\": 23,\\n        \\\"max\\\": 35,\\n
    \\\"num_unique_values\\\": 6,\\n        \\\"samples\\\": [\\n          25,\\n
    30,\\n          28\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n
    \\\"description\\\": \\\"\\\"\\n    }\\n    },\\n    {\\n      \\\"column\\\":
    \\\"Marks\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n
    \\\"std\\\": 6,\\n        \\\"min\\\": 75,\\n        \\\"max\\\": 95,\\n
    \\\"num_unique_values\\\": 6,\\n        \\\"samples\\\": [\\n          85,\\n
    90,\\n          95\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n
    \\\"description\\\": \\\"\\\"\\n    }\\n    },\\n    {\\n      \\\"column\\\":
    \\\"Department\\\",\\n      \\\"properties\\\": {\\n        \\\"dtype\\\":
    \\\"string\\\",\\n        \\\"num_unique_values\\\": 4,\\n        \\\"samples\\\":
    [\\n          \\\"Marketing\\\",\\n          \\\"Management\\\",\\n
    \\\"Design\\\"\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n
    \\\"description\\\": \\\"\\\"\\n    }\\n    }\\n  ]\\n}\\",\"type\":\"dataframe\"}

```

Missing Data

```

data = {
    "Name": ["Alice", "Bob", "Charlie", "Diana", "Eve", "Frank",
"Grace", "Hank"],
    "Gender": ["Female", "Male", None, "Female", None, "Male",
"Female", None],
}
df = pd.DataFrame(data)
print(df)

```

	Name	Gender
0	Alice	Female
1	Bob	Male
2	Charlie	None
3	Diana	Female
4	Eve	None
5	Frank	Male
6	Grace	Female
7	Hank	None

```
print("\nCheck for missing values:")
print(pd.isnull(df))
```

Check for missing values:

	Name	Gender
0	False	False
1	False	False
2	False	True
3	False	False
4	False	True
5	False	False
6	False	False
7	False	True

```
print("\nCheck for missing values(Column):")
print(pd.isnull(df['Gender']))
```

Check for missing values(Column):

0	False
1	False
2	True
3	False
4	True
5	False
6	False
7	True

Name: Gender, dtype: bool

```
# Fill missing values in the 'Gender' column with a default value
df['Gender'] = df['Gender'].fillna("Not Specified")
```

#updated dataframe

```
print(df)
```

	Name	Gender
0	Alice	Female
1	Bob	Male
2	Charlie	Not Specified
3	Diana	Female
4	Eve	Not Specified
5	Frank	Male
6	Grace	Female
7	Hank	Not Specified

#Read myStudentDataFrame.csv

```
import pandas as pd
df = pd.read_csv('myDataFrame_12340220.csv')
print(df)
```


	Name	Age	Marks	Department
0	Alice	25	85	Design
1	Bob	30	90	Marketing
2	Charlie	35	88	Development
3	James	23	75	Management
4	John	26	80	Marketing
5	William	28	95	Design
6	Caleb	25	85	Development
7	Helen	30	90	Design

```
#Check for missing data in all columns using appropriate pandas functions.
```

```
missing_data = df.isnull().sum()
print("Missing Data in Columns:\n", missing_data)
```

```
Missing Data in Columns:
```

```
  Name      0
Age      0
Marks     0
Department 0
dtype: int64
```

```
#Drop Rows with Missing Data
```

```
df_cleaned = df.dropna()
```

```
#Compute Summary Statistics (AVG,MEAN,MAX,MIN)
```

```
summary_statistics = {
    "Mean": df_cleaned['Marks'].mean(),
    "Max": df_cleaned['Marks'].max(),
    "Min": df_cleaned['Marks'].min(),
    "Avg": df_cleaned['Marks'].mean()
}
print("\nSummary Statistics:\n", summary_statistics)
```

```
Summary Statistics:
```

```
{'Mean': 86.0, 'Max': 95, 'Min': 75, 'Avg': 86.0}
```

```
#Filter Data and Compute Pass/Fail
```

```
#mark >= 40: Pass
```

```
#mark < 40: Fail
```

```
#Add a new column Result to the DataFrame indicating Pass or Fail.
```

```
df_cleaned['Result'] = df_cleaned['Marks'].apply(lambda x: 'Pass' if x
>= 40 else 'Fail')
```

```
#Save the Final DataFrame
```

```
#Save the updated DataFrame (with the Result column) to a new CSV file
named Result_YourIDNumber.csv.
```

```
df_cleaned.to_csv('Result_12340220.csv', index=False)
```

```
print("\nUpdated DataFrame with 'Result' column saved to  
'Result_12340220.csv'.")
```

Updated DataFrame with 'Result' column saved to 'Result_12340220.csv'.

Additional Practice Questions - [Click Here](#)

